AI Agent

Xu Tan/谭旭

Outline

- Background
  - Agent & AGI
  - What is Agent
- Agent: Foundations
  - Key Components: Reasoning, Memory, Tool Use
  - Agentic Workflow vs Large Agent Model
- Agent: Applications
  - Search/Research Agent
  - Computer-Using Agent
  - Other Vertical Agents
- Challenges & Future Trends

AGI Roadmap


OpenAI's 5-Level Plan to AGI
Conversational AI
1
Reasoners
2
Agents
3
Innovators
4
Organizations
5
chatgpt4online.org

AGI Roadmap: From Chatbot to Reasoner to Agent

Level 1: Chatbot
(Language model)

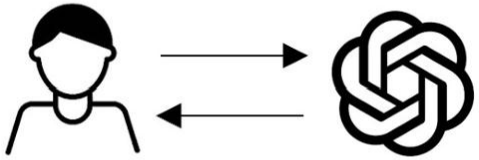

direct respond

Level 2: Reasoner
(Reasoning model)

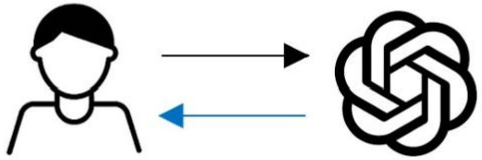

slow think
before respond

Level 3: Agent
(Agent model)

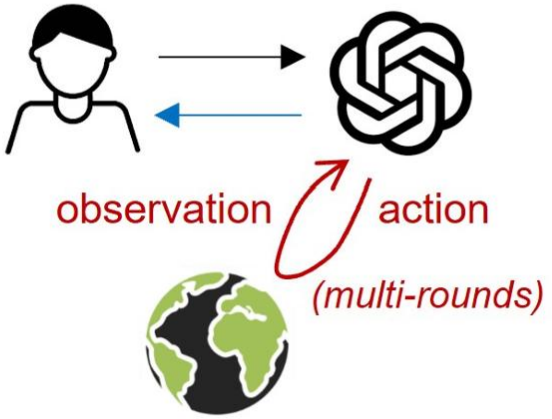

iterative slow think & action
before respond

From 2503.06580

What is AI Agent

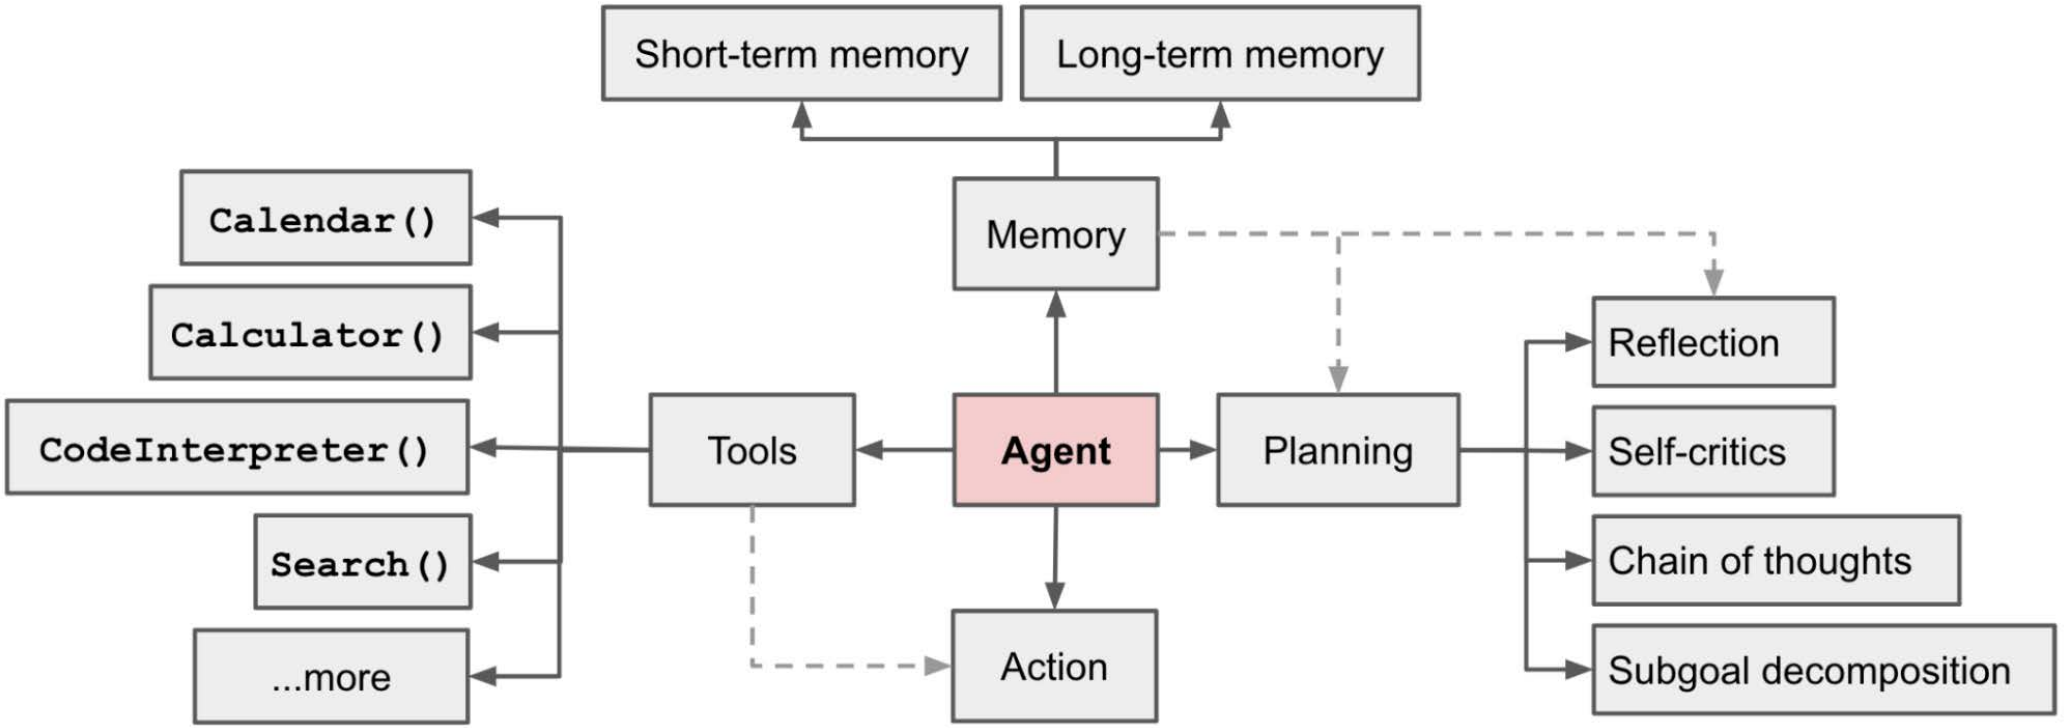

From Lilian Weng's blog

Outline

- Background
  - Agent & AGI
  - What is Agent
- Agent: Foundations
  - Key Components: Reasoning, Memory, Tool Use
  - Agentic Workflow vs Large Agent Model
- Agent: Applications
  - Search/Research Agent
  - Computer-Using Agent
  - Other Vertical Agents
- Challenges & Future Trends

Key Components of AI Agent

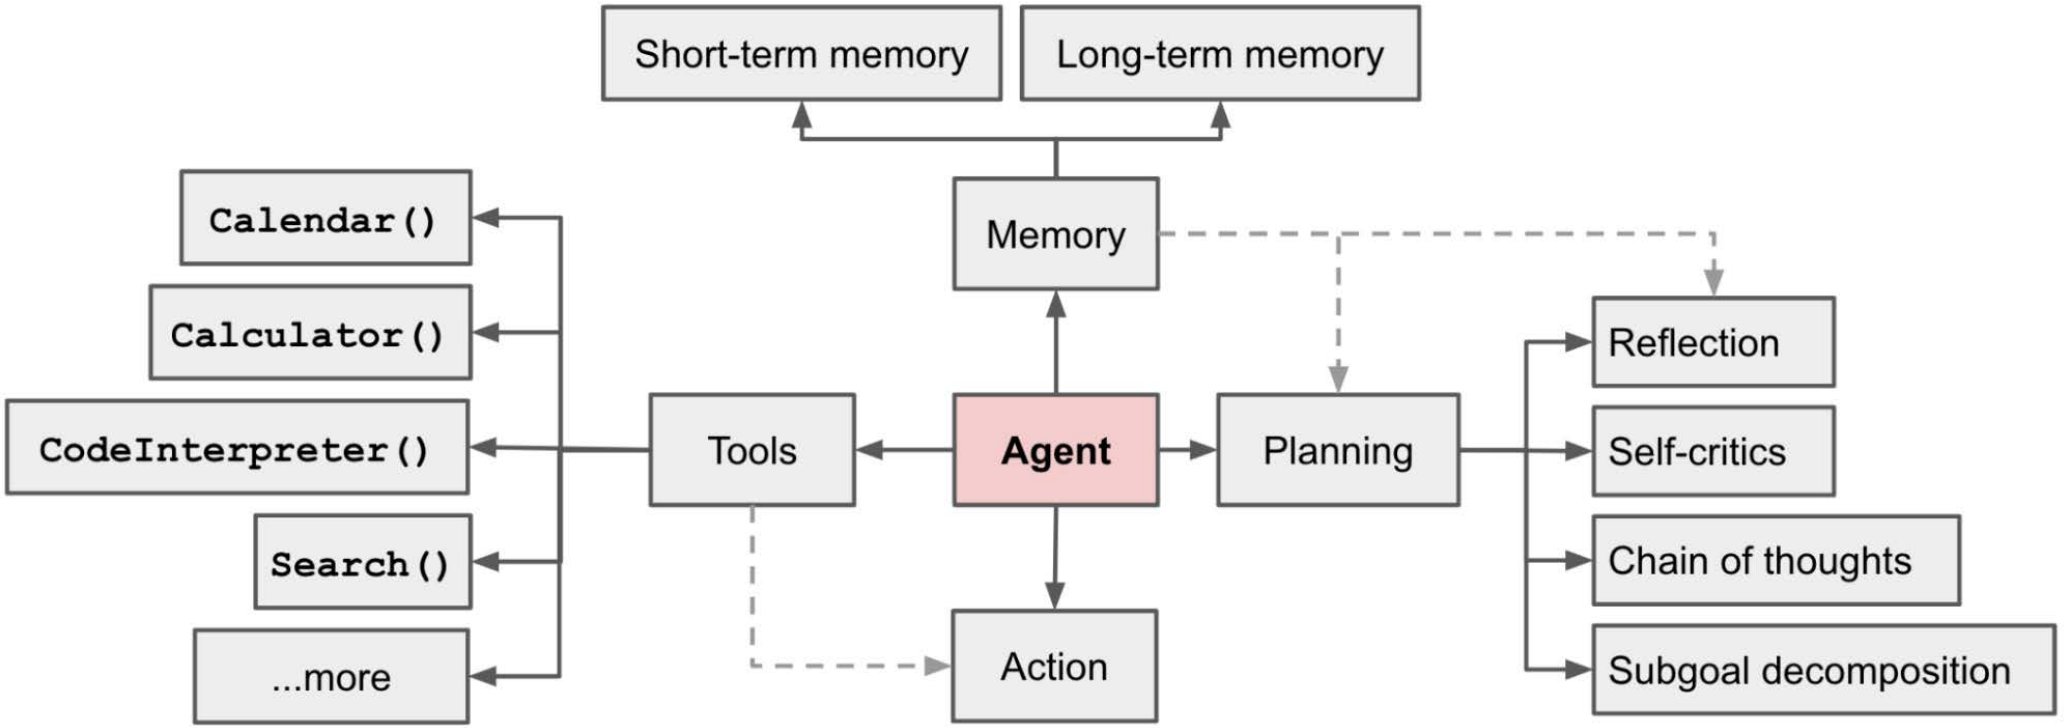

From Lilian Weng's blog

Progress on Reasoning/Memory/Tool

From 硅谷101

Key Components of AI Agent: Reasoning

- Planning, task decomposition, chain-of-thought, reflection, reasoning
  - Multi-step reasoning: CoT step-by-step, search tree
  - Self-refection: ReAct: Synergizing reasoning and acting; Reflexion: reinforce language agents by linguistic feedback
  - Intrinsic Reasoning: OpenAI O1 & DeepSeek R1

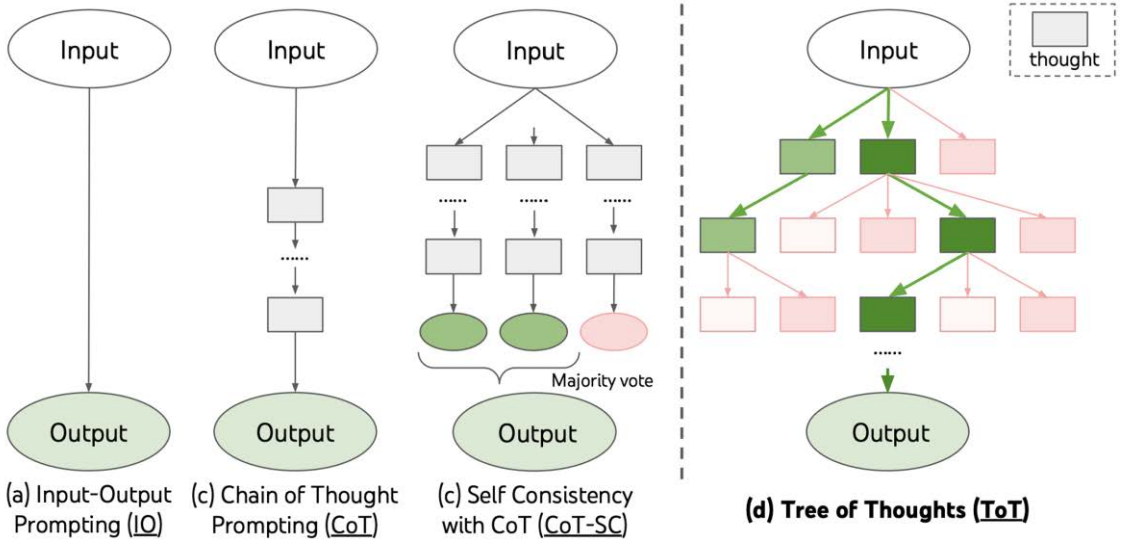

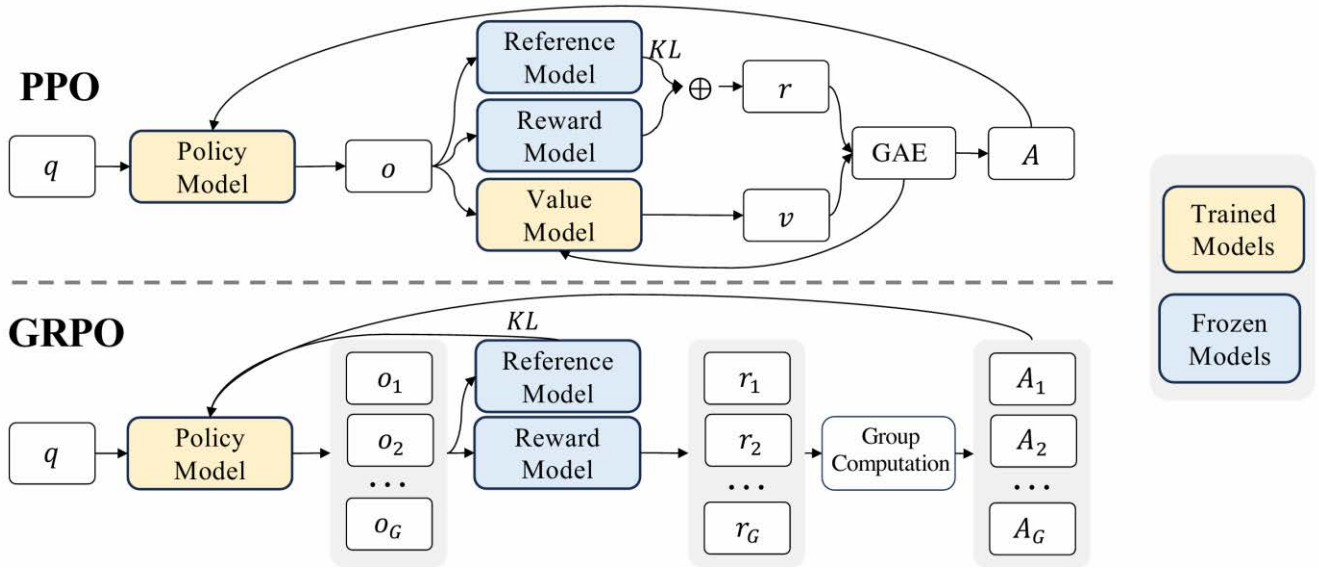

From 2305.10601 2402.03300

Key Components of AI Agent: Memory

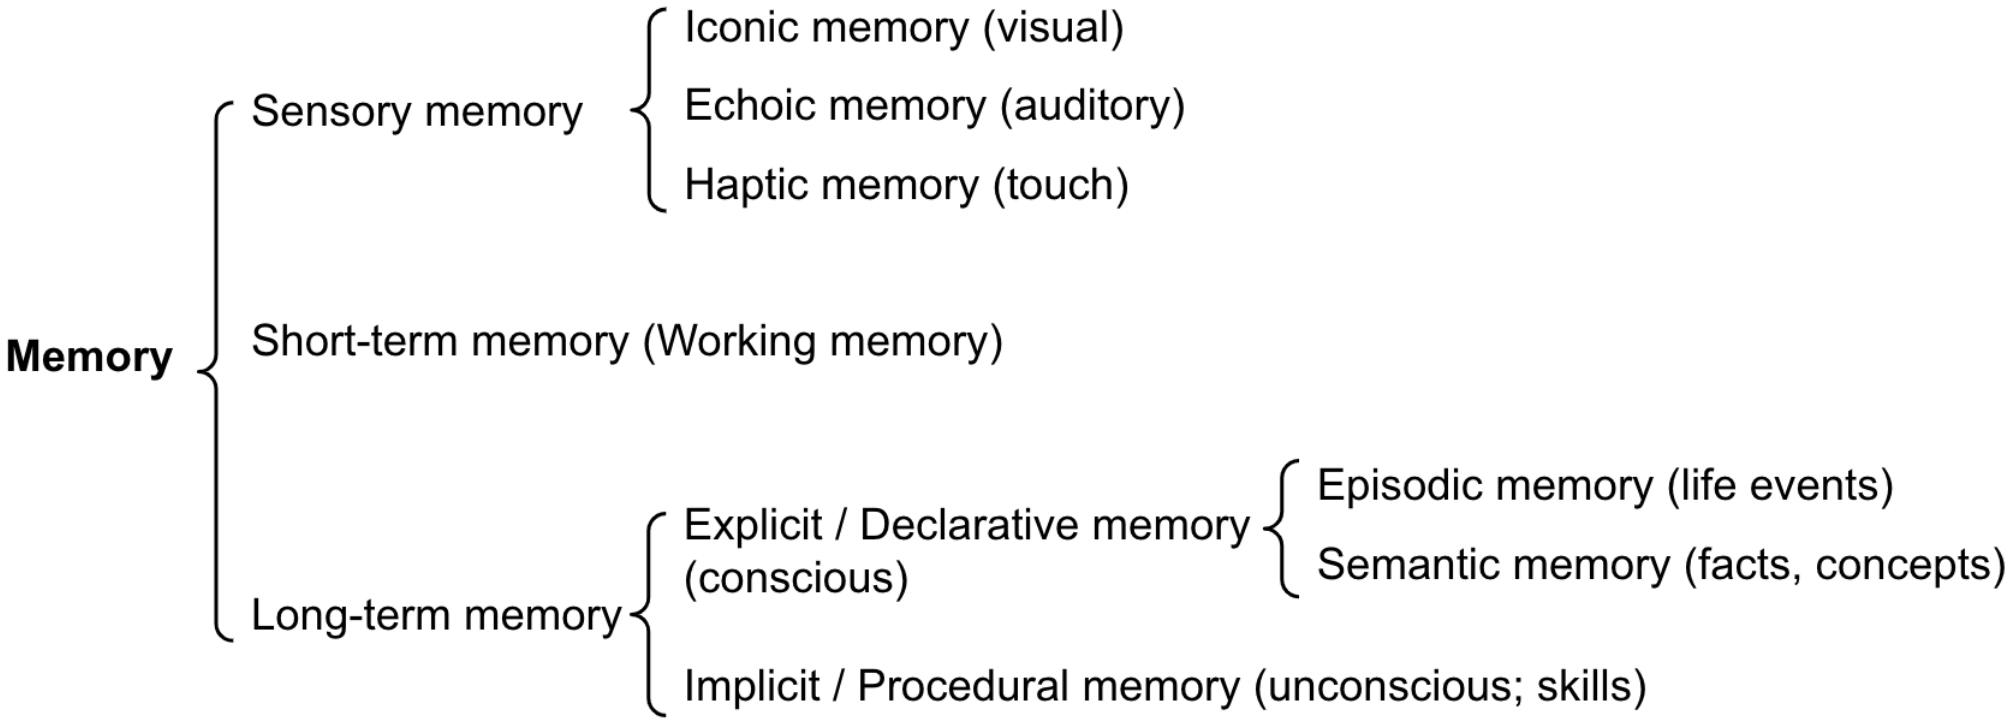

- Sensory memory: embedding input
- Short-term memory: in-context learning
- Long-term memory: retrieval-augmented generation (RAG)

From Lilian Weng's blog

Key Components of AI Agent: Memory

- Progress on Long-Context

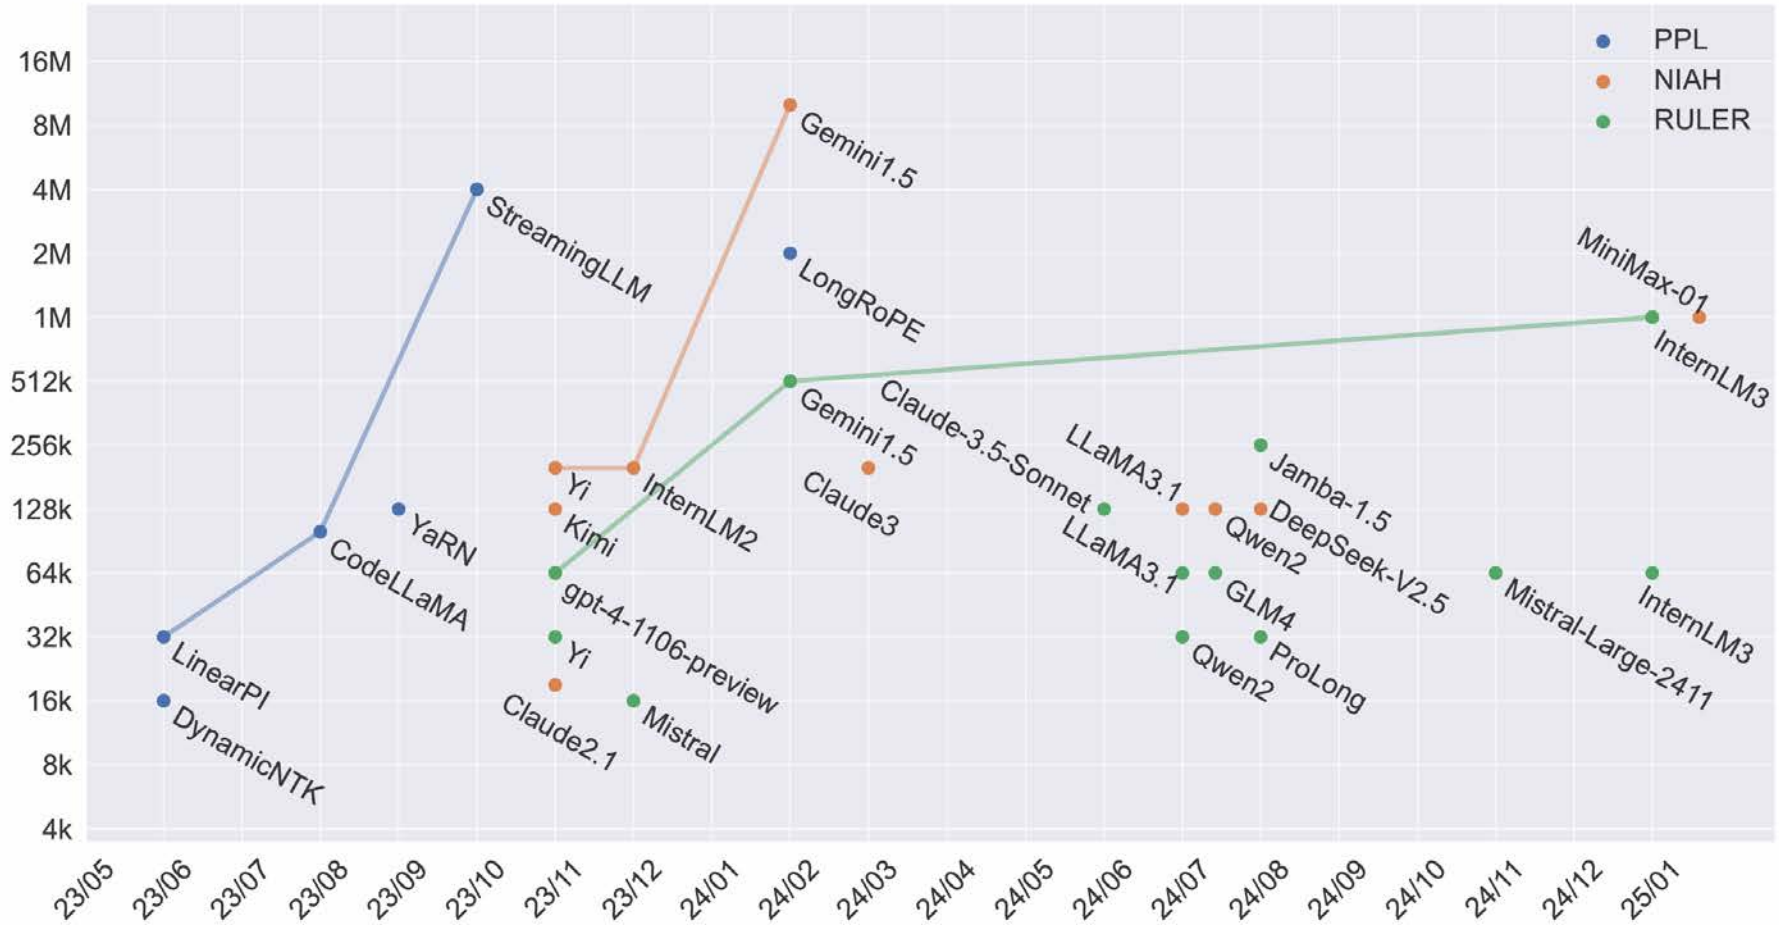

From 2502.17129

Key Components of AI Agent: Memory

- Storage: Reduce KV cache memory

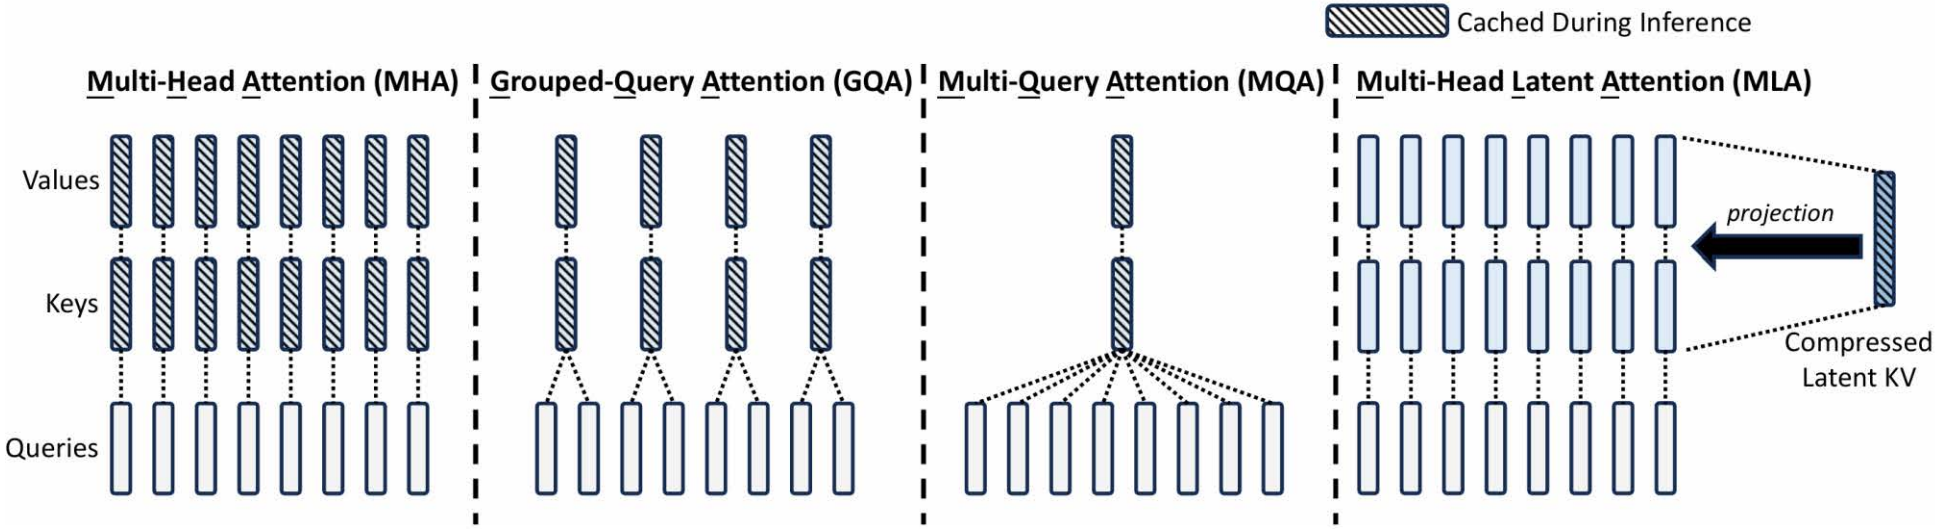

- Computation: Sparse/linear attention
  - Native sparse attention, mixture of block attention
  - Mamba

From 2405.04434 2502.11089 2502.13189 2312.00752

Key Components of AI Agent: Tool

- Toolformer

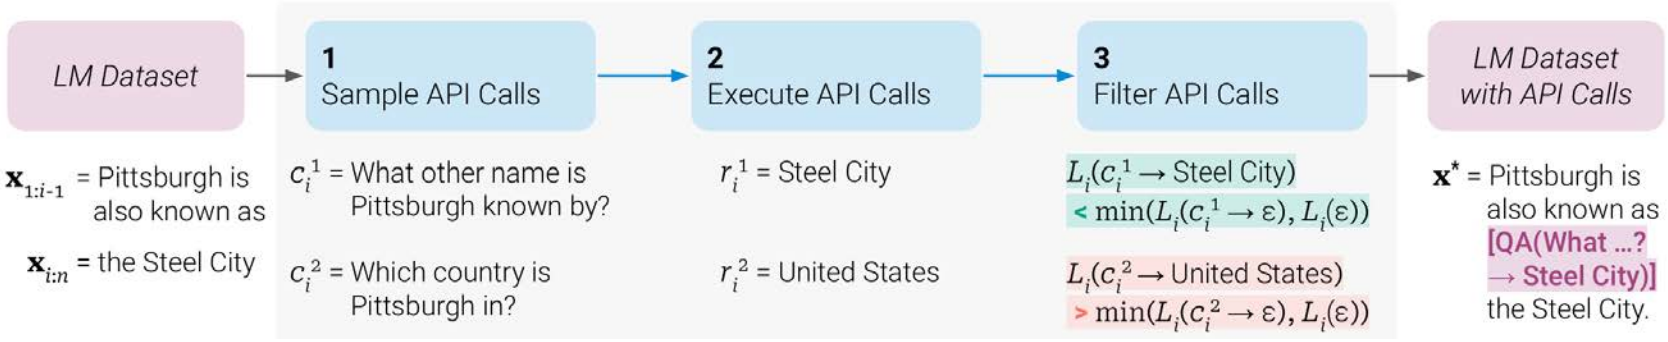

The New England Journal of Medicine is a registered trademark of [QA("Who is the publisher of The New England Journal of Medicine?") → Massachusetts Medical Society] the MMS.

Out of 1400 participants, 400 (or [Calculator(400 / 1400) → 0.29] 29%) passed the test.

The name derives from "la tortuga", the Spanish word for [MT("tortuga") → turtle] turtle.

The Brown Act is California's law [WikiSearch("Brown Act") → The Ralph M. Brown Act is an act of the California State Legislature that guarantees the public's right to attend and participate in meetings of local legislative bodies.] that requires legislative bodies, like city councils, to hold their meetings open to the public.

From 2302.04761

Key Components of AI Agent: Tool

- HuggingGPT
  - Task planning, model selection, task execution, response generation

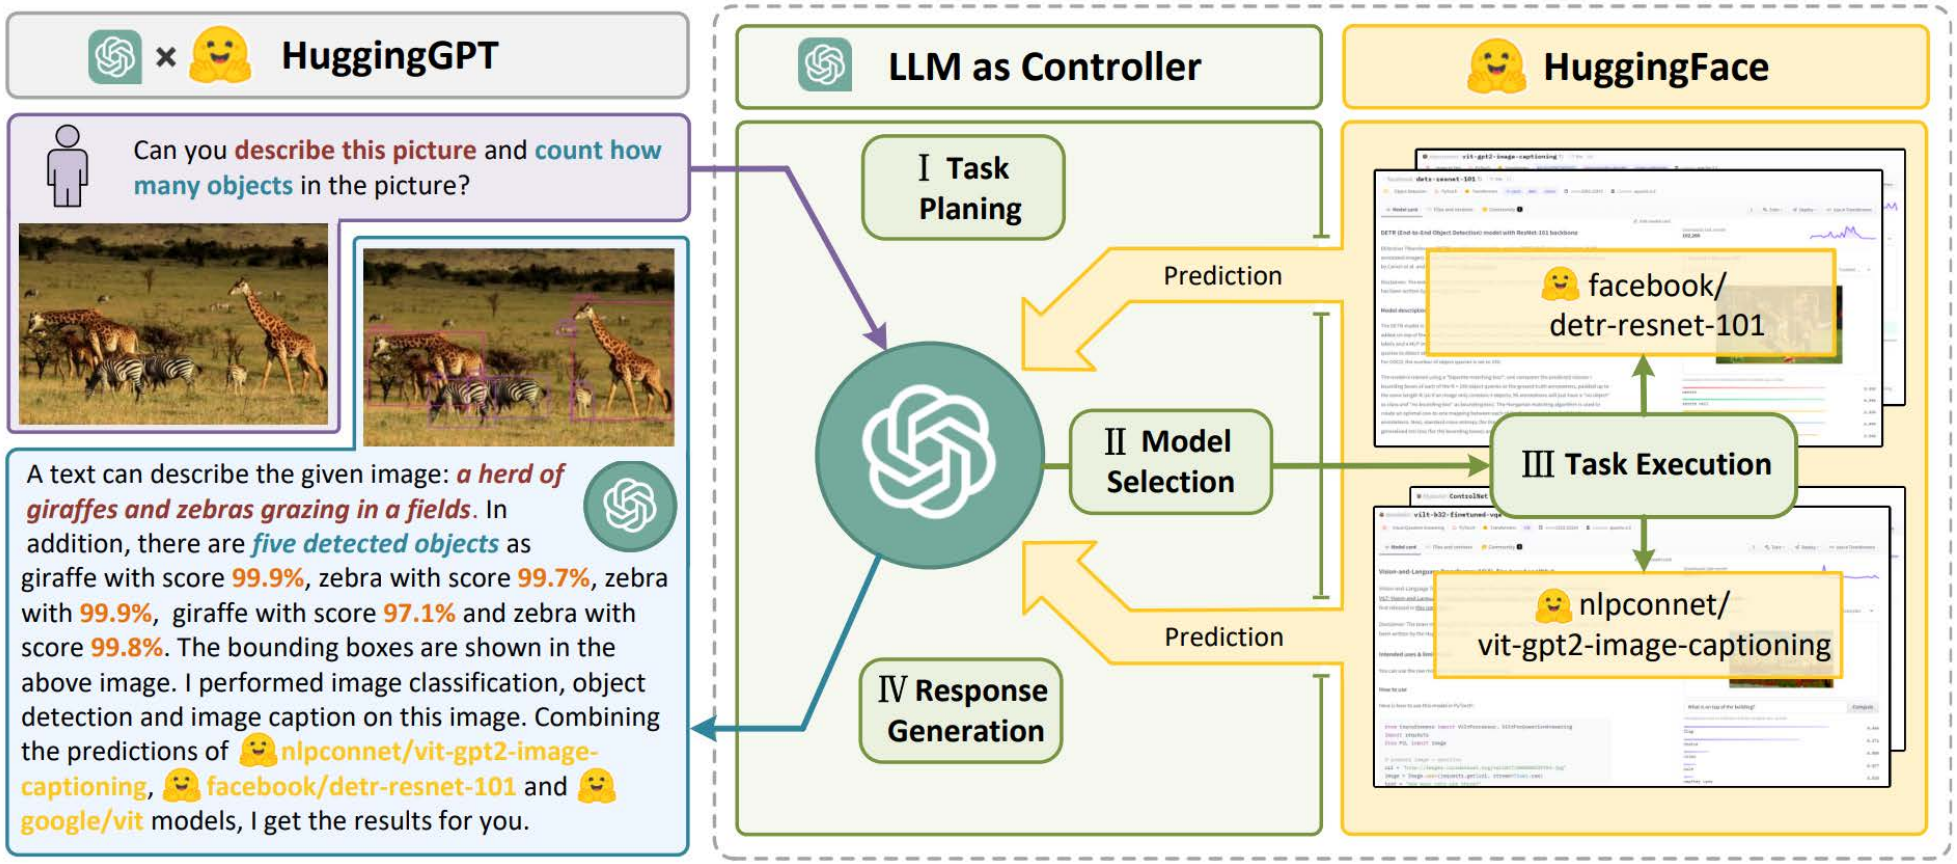

From 2303.17580

Key Components of AI Agent: Tool

- Browsers and computers
  - Use browsers and computers through API or GUI

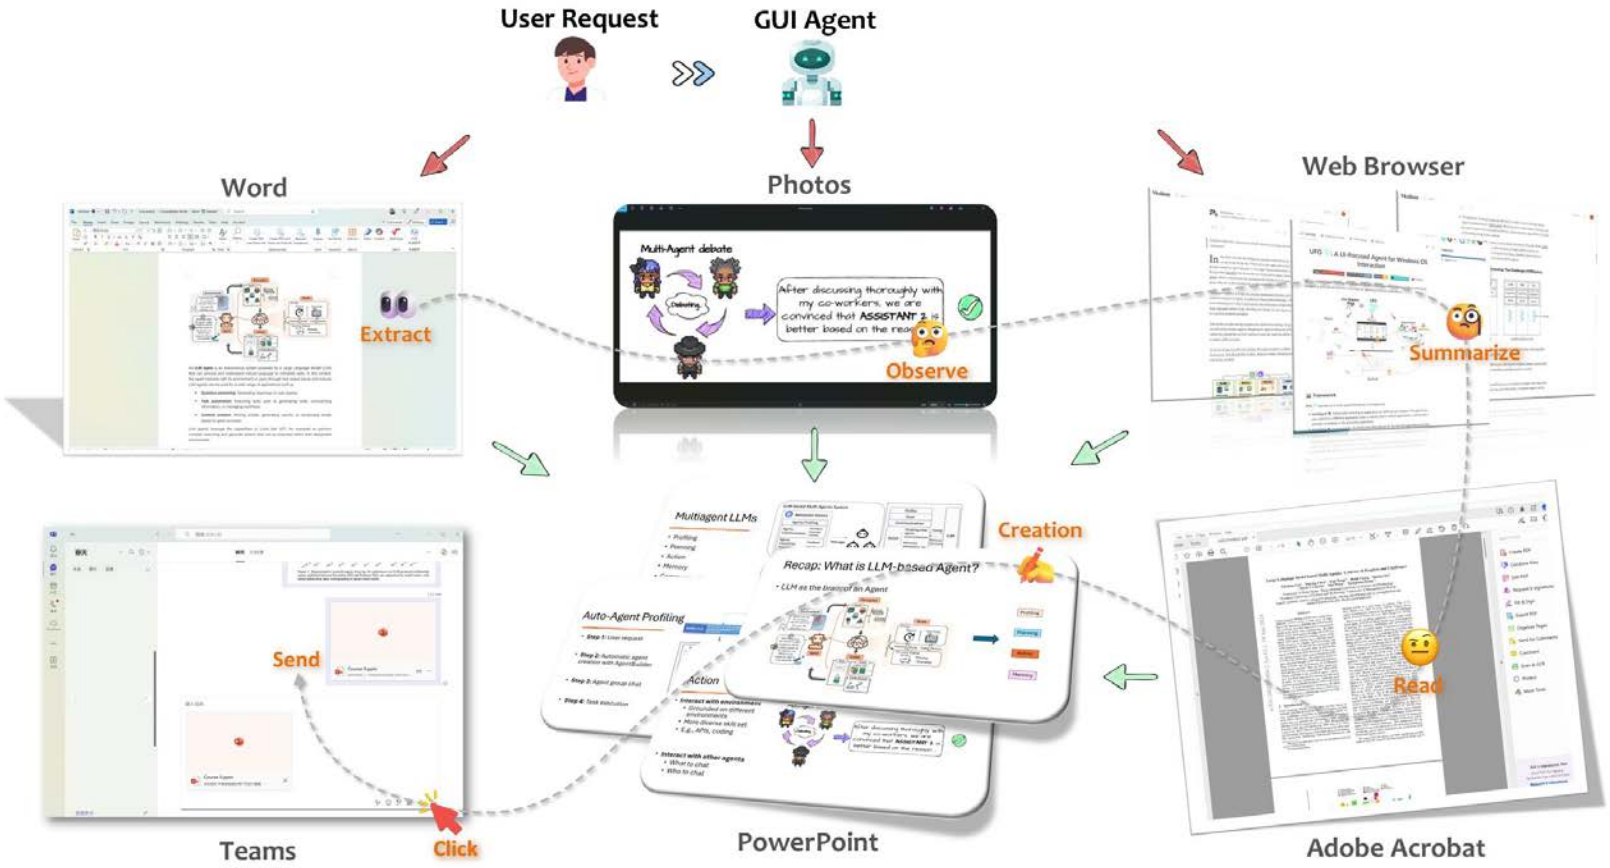

From 2411.18279

Key Components of AI Agent: Tool

- Model Context Protocol (MCP)

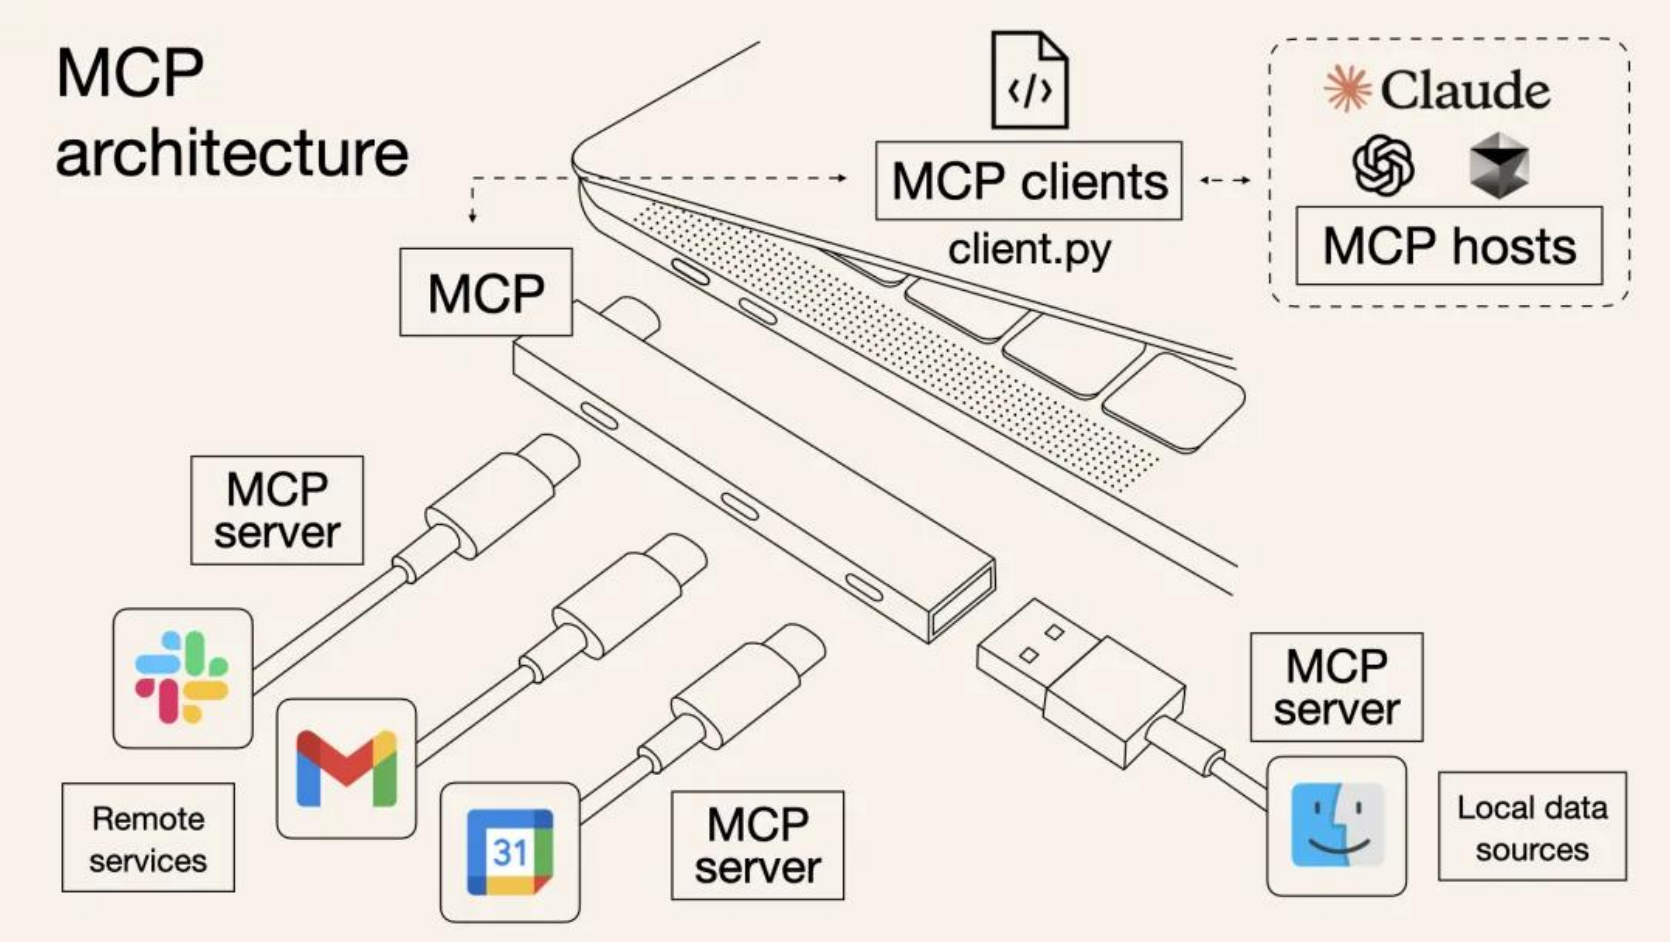

Key Components of AI Agent: Tool

- Model Context Protocol (MCP)

1. 客户端（Claude Desktop / Cursor）将你的问题发送给 Claude。
2. Claude 分析可用的工具，并决定使用哪一个（或多个）。
3. 客户端通过 MCP Server 执行所选的工具。
4. 工具的执行结果被送回给 Claude。
5. Claude 结合执行结果构造最终的 prompt 并生成自然语言的回应。
6. 回应最终展示给用户！

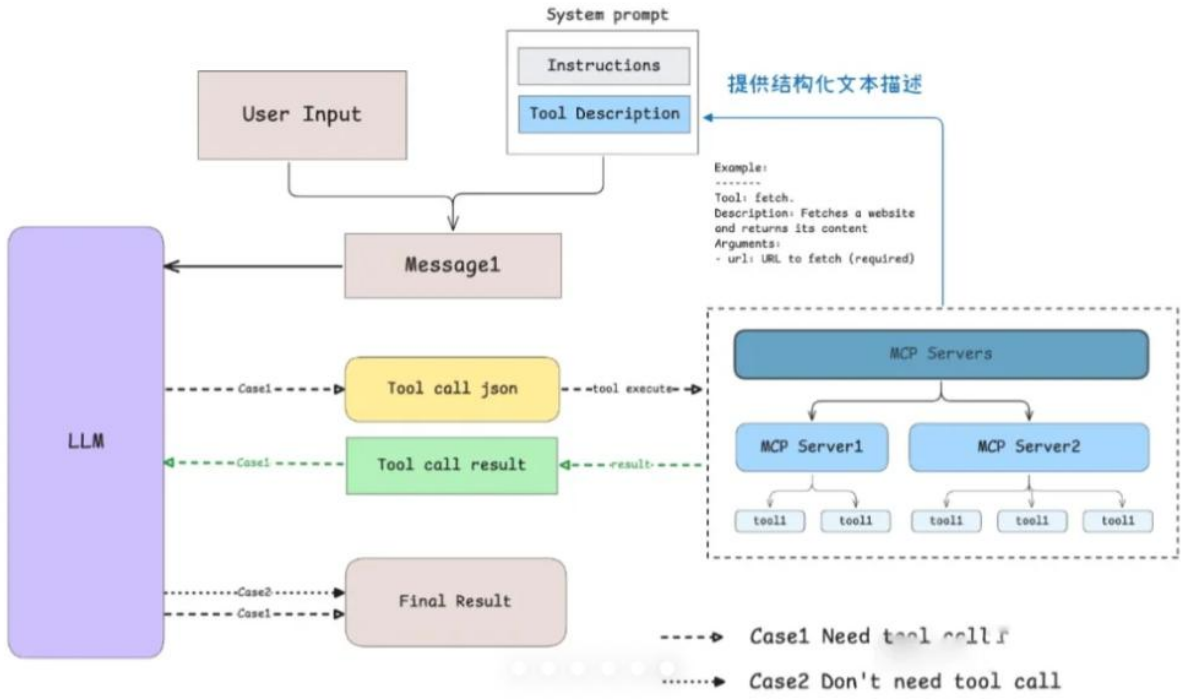

Key Components of AI Agent: Tool

- Model Context Protocol (MCP)
  - An example: LLM call Ableton to create music
    - LLM: Claude
    - LLM Client: Claude Desktop
    - MCP server: AbletonMCP Server
      - AbletonMCP server controls Ableton through API

From ahujasid/ableton-mcp

Outline

- Background
  - Agent & AGI
  - What is Agent
- Agent: Foundations
  - Key Components: Reasoning, Memory, Tool Use
  - Agentic Workflow vs Large Agent Model
- Agent: Applications
  - Search/Research Agent
  - Computer-Using Agent
  - Other Vertical Agents
- Challenges & Future Trends

Agentic Workflow vs Large Agent Model

- Agentic workflow
  - Prompting-based, predefined workflow

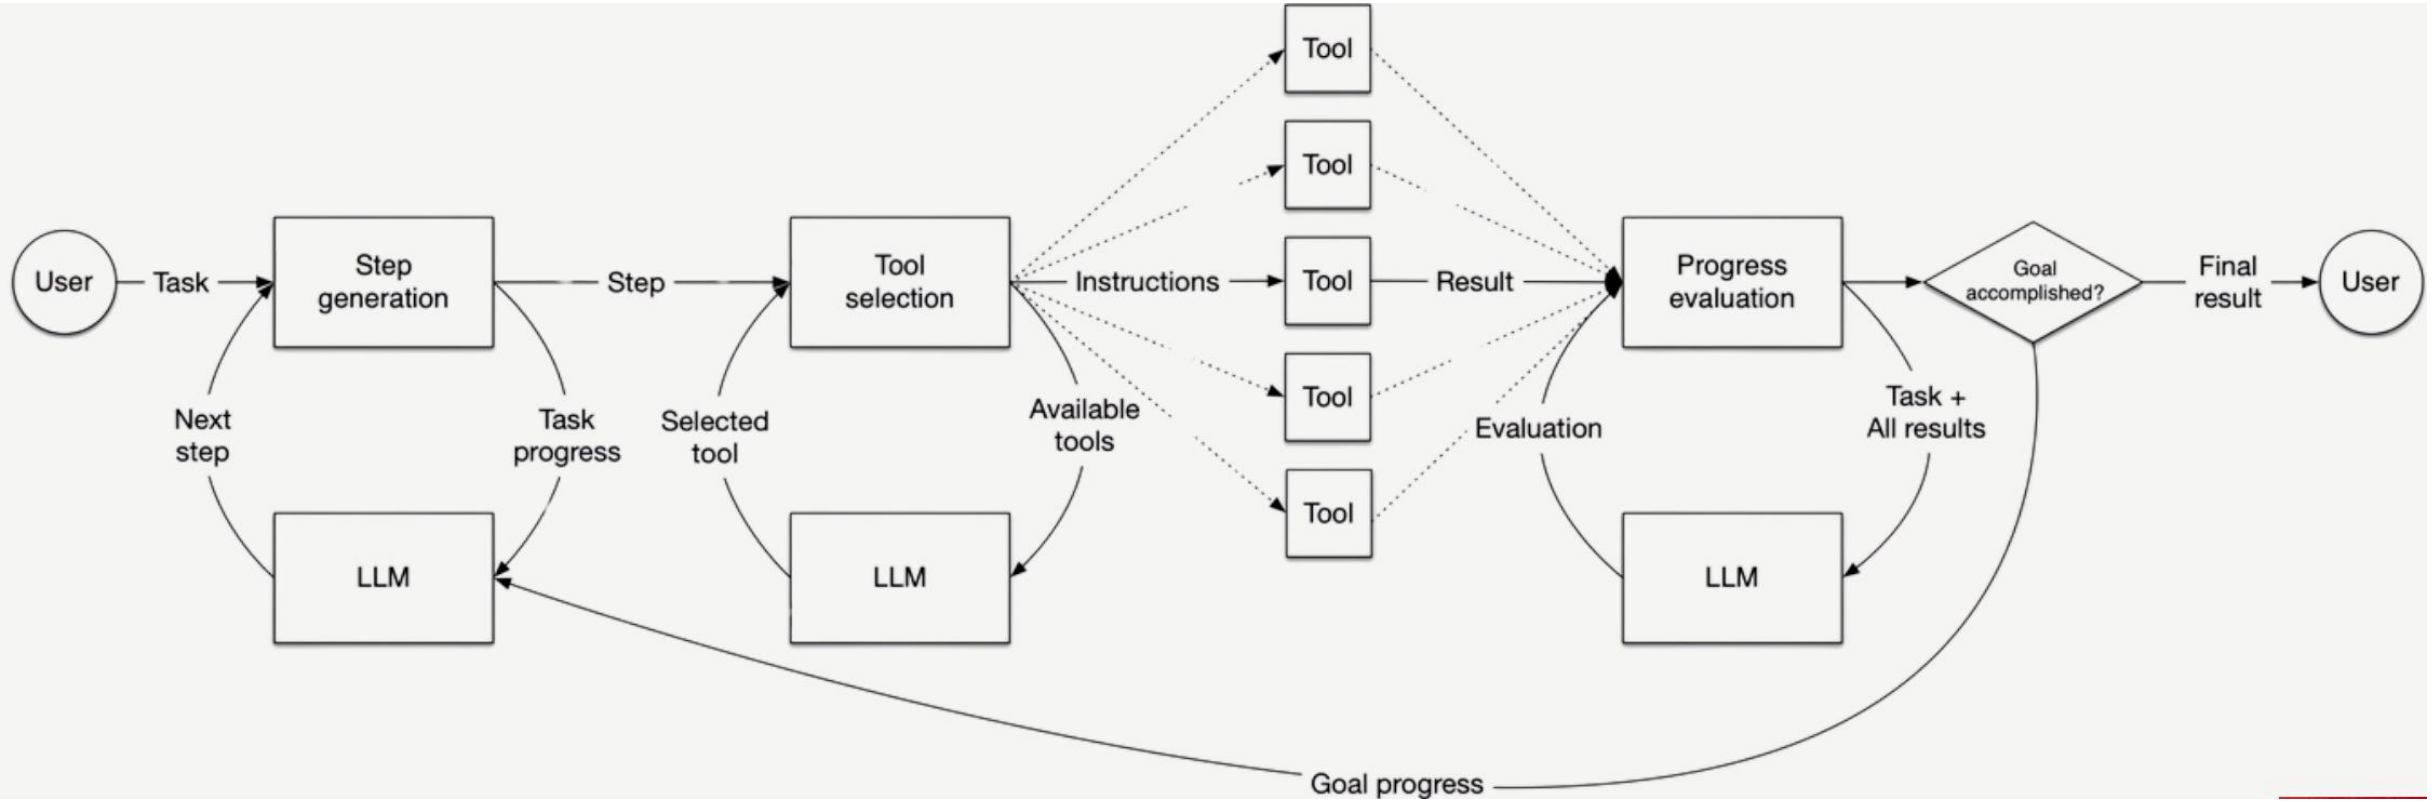

Agentic Workflow vs Large Agent Model

- From prompting to learning

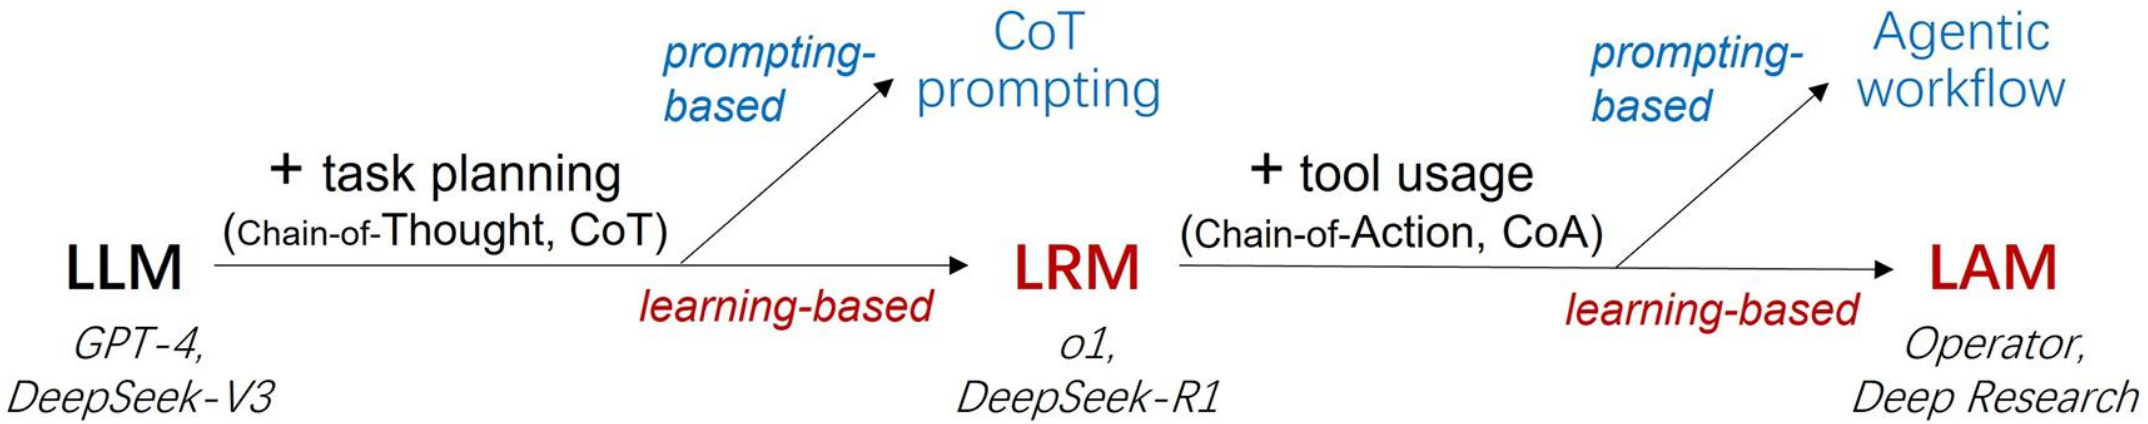

From 2503.06580

Agentic Workflow vs Large Agent Model

- Less structure, more intelligence

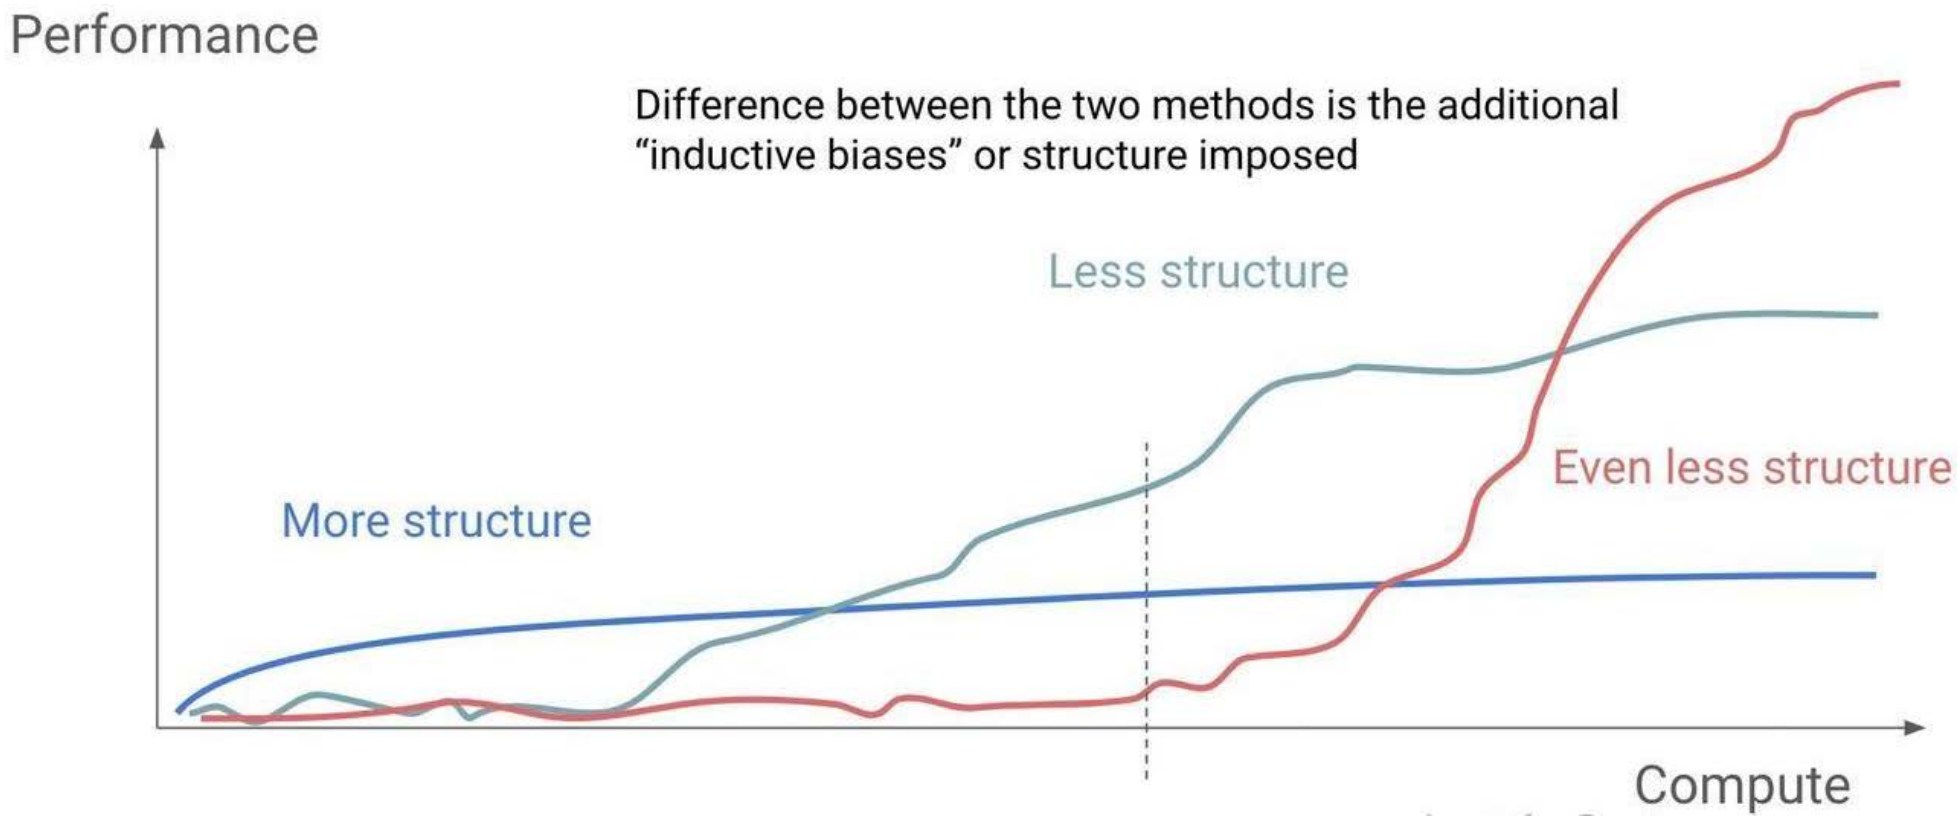

Agentic Workflow vs Large Agent Model

- Large agent model (LAM)
  - Internalize: learn the logics between thinking steps and internalize the CoT generation capability as a “active” model behavior
  - Through SFT and/or RL
    - Data/environment preparation
      - e.g., simulate web search environment
      - <question><think><action><observation>…<think><action><observation><answer>
    - SFT: Warm start the model
    - Prompt and verifier/reward
      - How to speedup rollout process
      - How to get verifier/reward
    - RL: PPO/GRPO
    - Repeat

Outline

- Background
  - Agent & AGI
  - What is Agent
- Agent: Foundations
  - Key Components: Reasoning, Memory, Tool Use
  - Agentic Workflow vs Large Agent Model
- Agent: Applications
  - Search/Research Agent
  - Computer-Using Agent
  - Other Vertical Agents
- Challenges & Future Trends

Agent Application: Search/Research Agent

- The roadmap of search/research agent

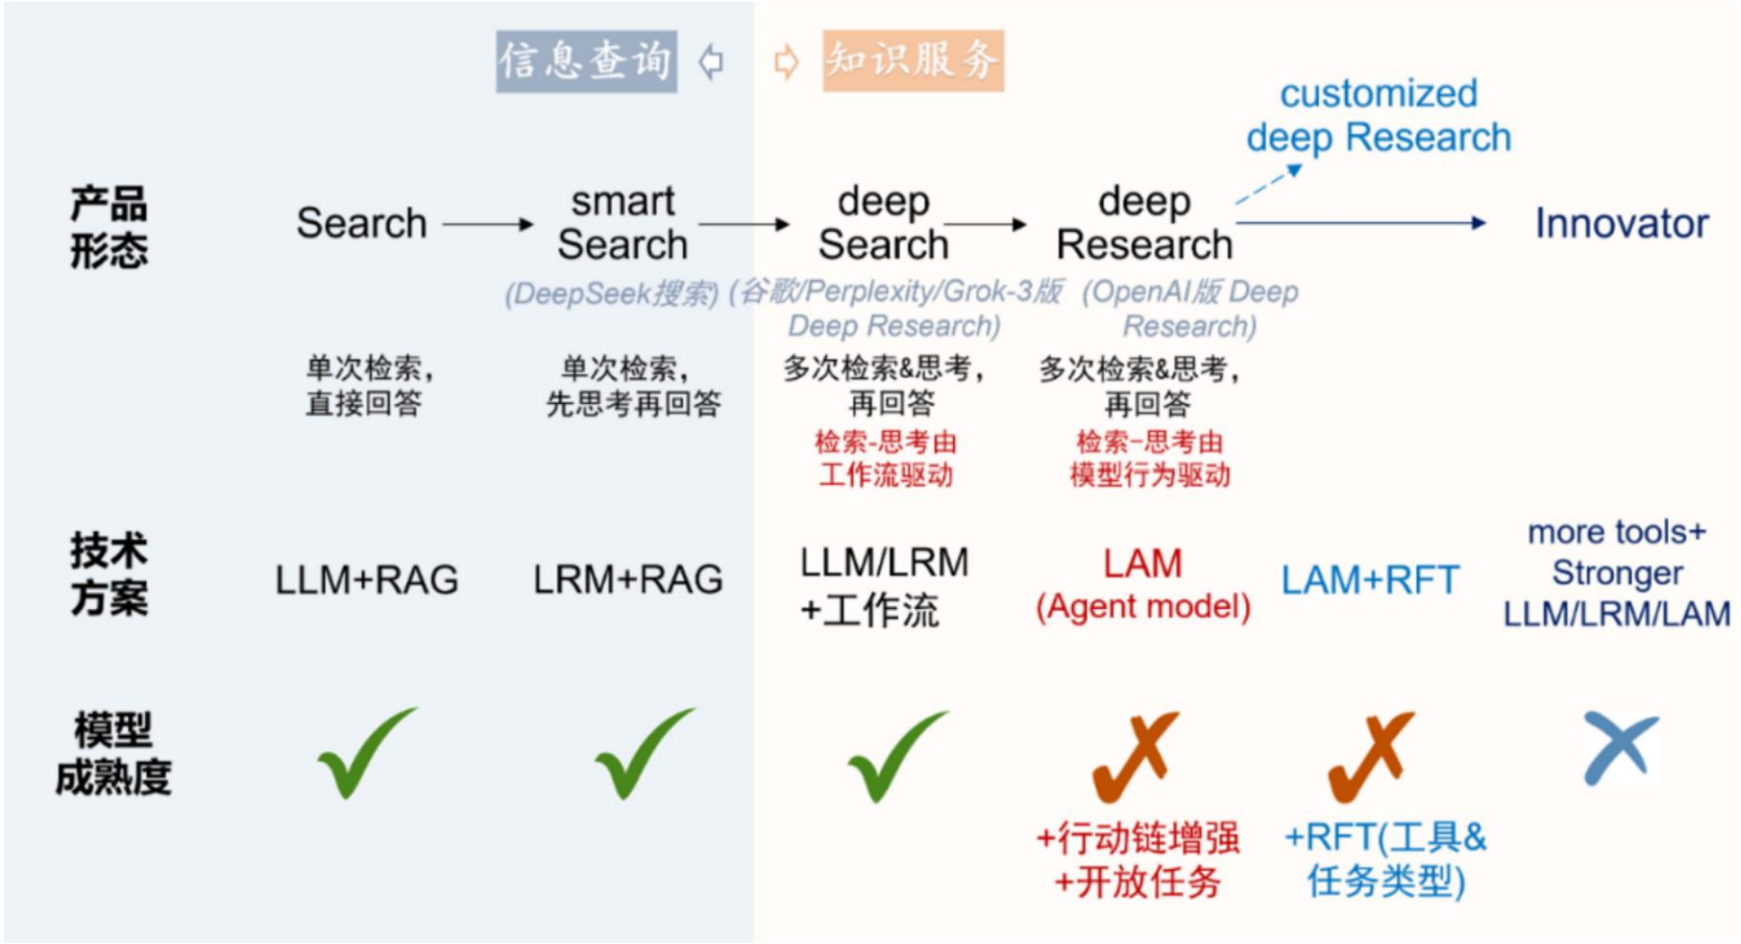

Agent Application: Computer-using Agent

- Computer-using Agent
  - Agents using computing devices (e.g., computers and mobile phones) by operating within the environments and interfaces (e.g., Graphical User Interface (GUI), Application Programming Interface (API)) provided by operating systems (OS) to automate tasks

  - Computing devices: Computers, mobile phones, servers
  - OS Agent (windows, mac, android, ios), Web Agent
  - API agent vs GUI Agent

(a) Web GUI (b) Mobile GUI (c) Computer GUI

Agent Application: Computer-using Agent

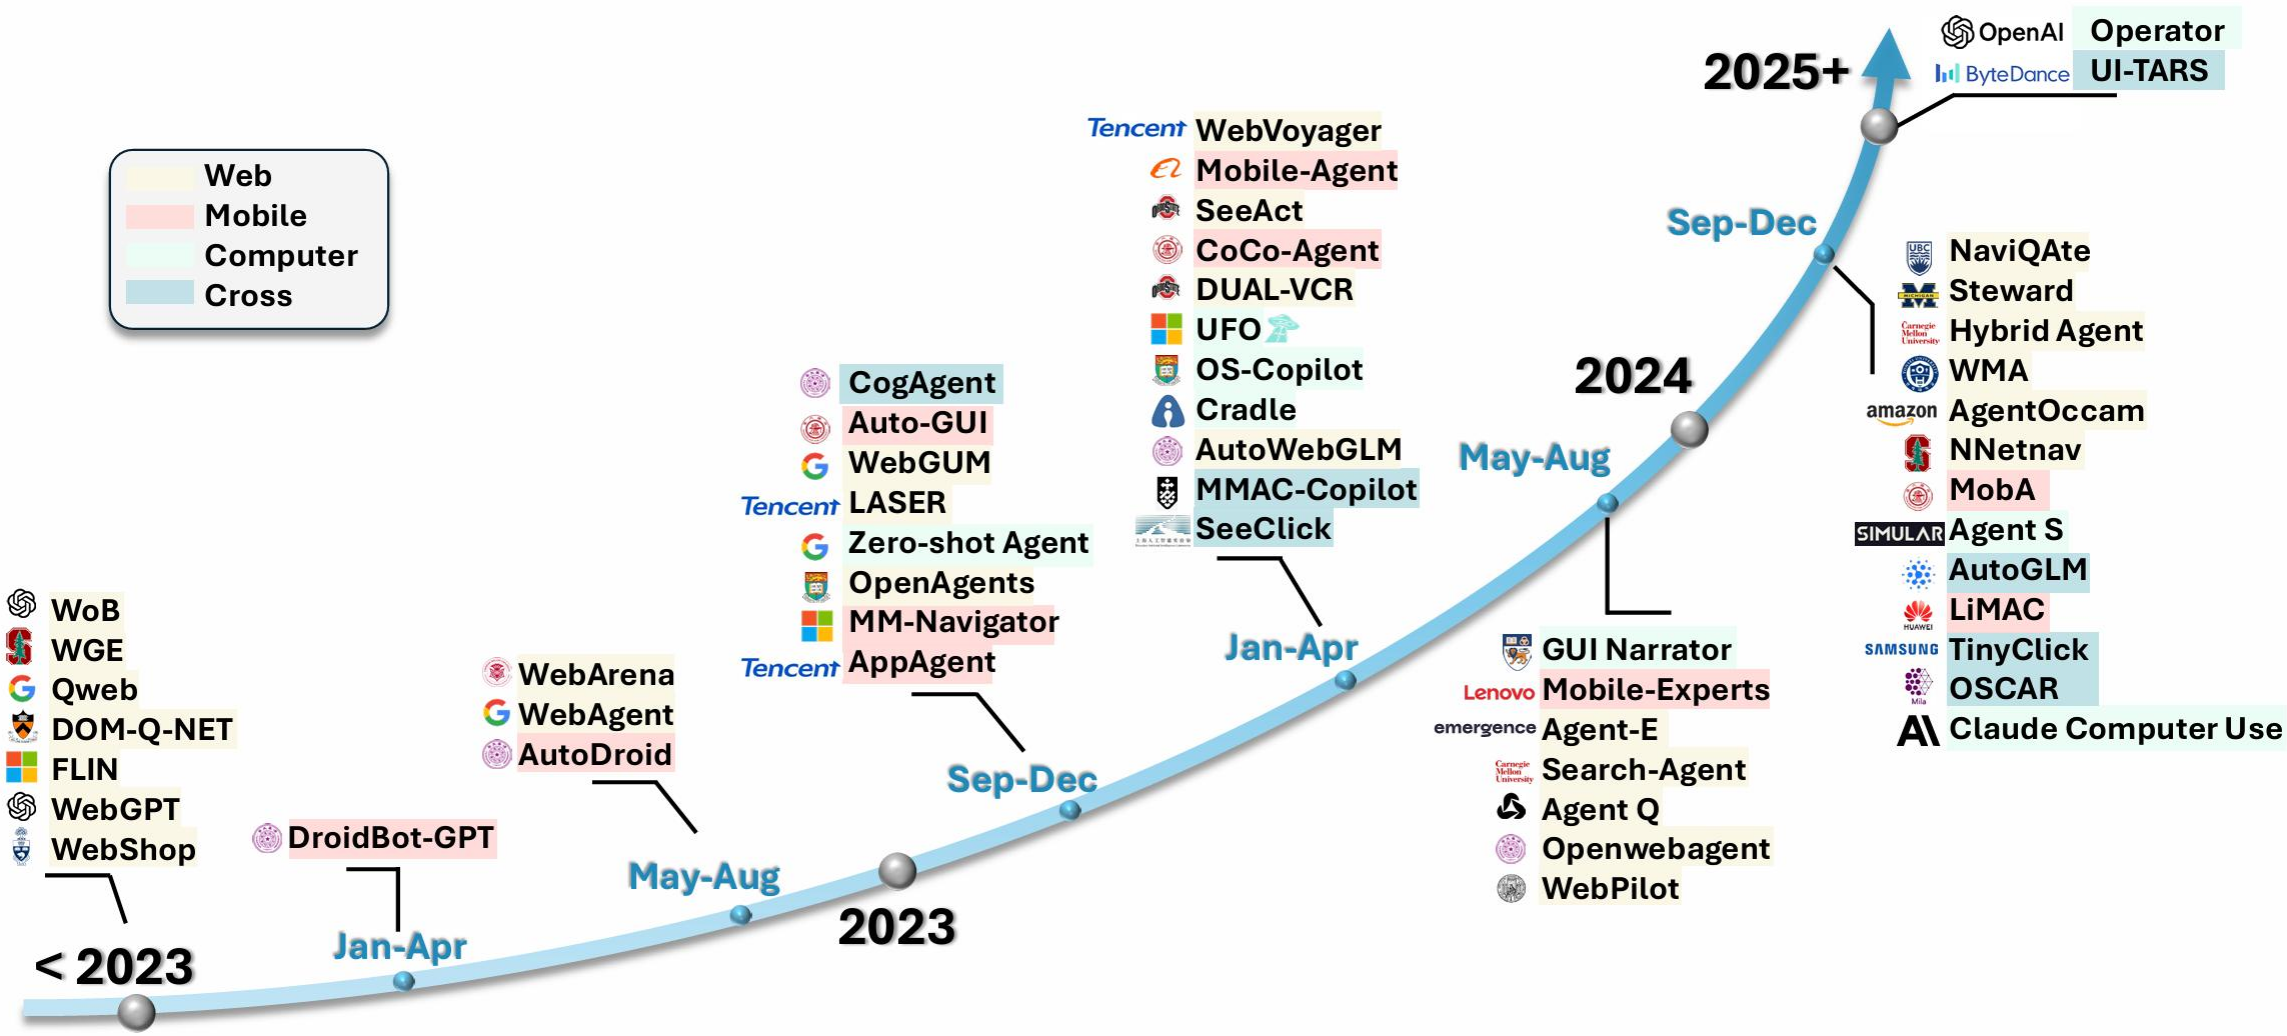

From 2411.18279

Agent Application: Computer-using Agent

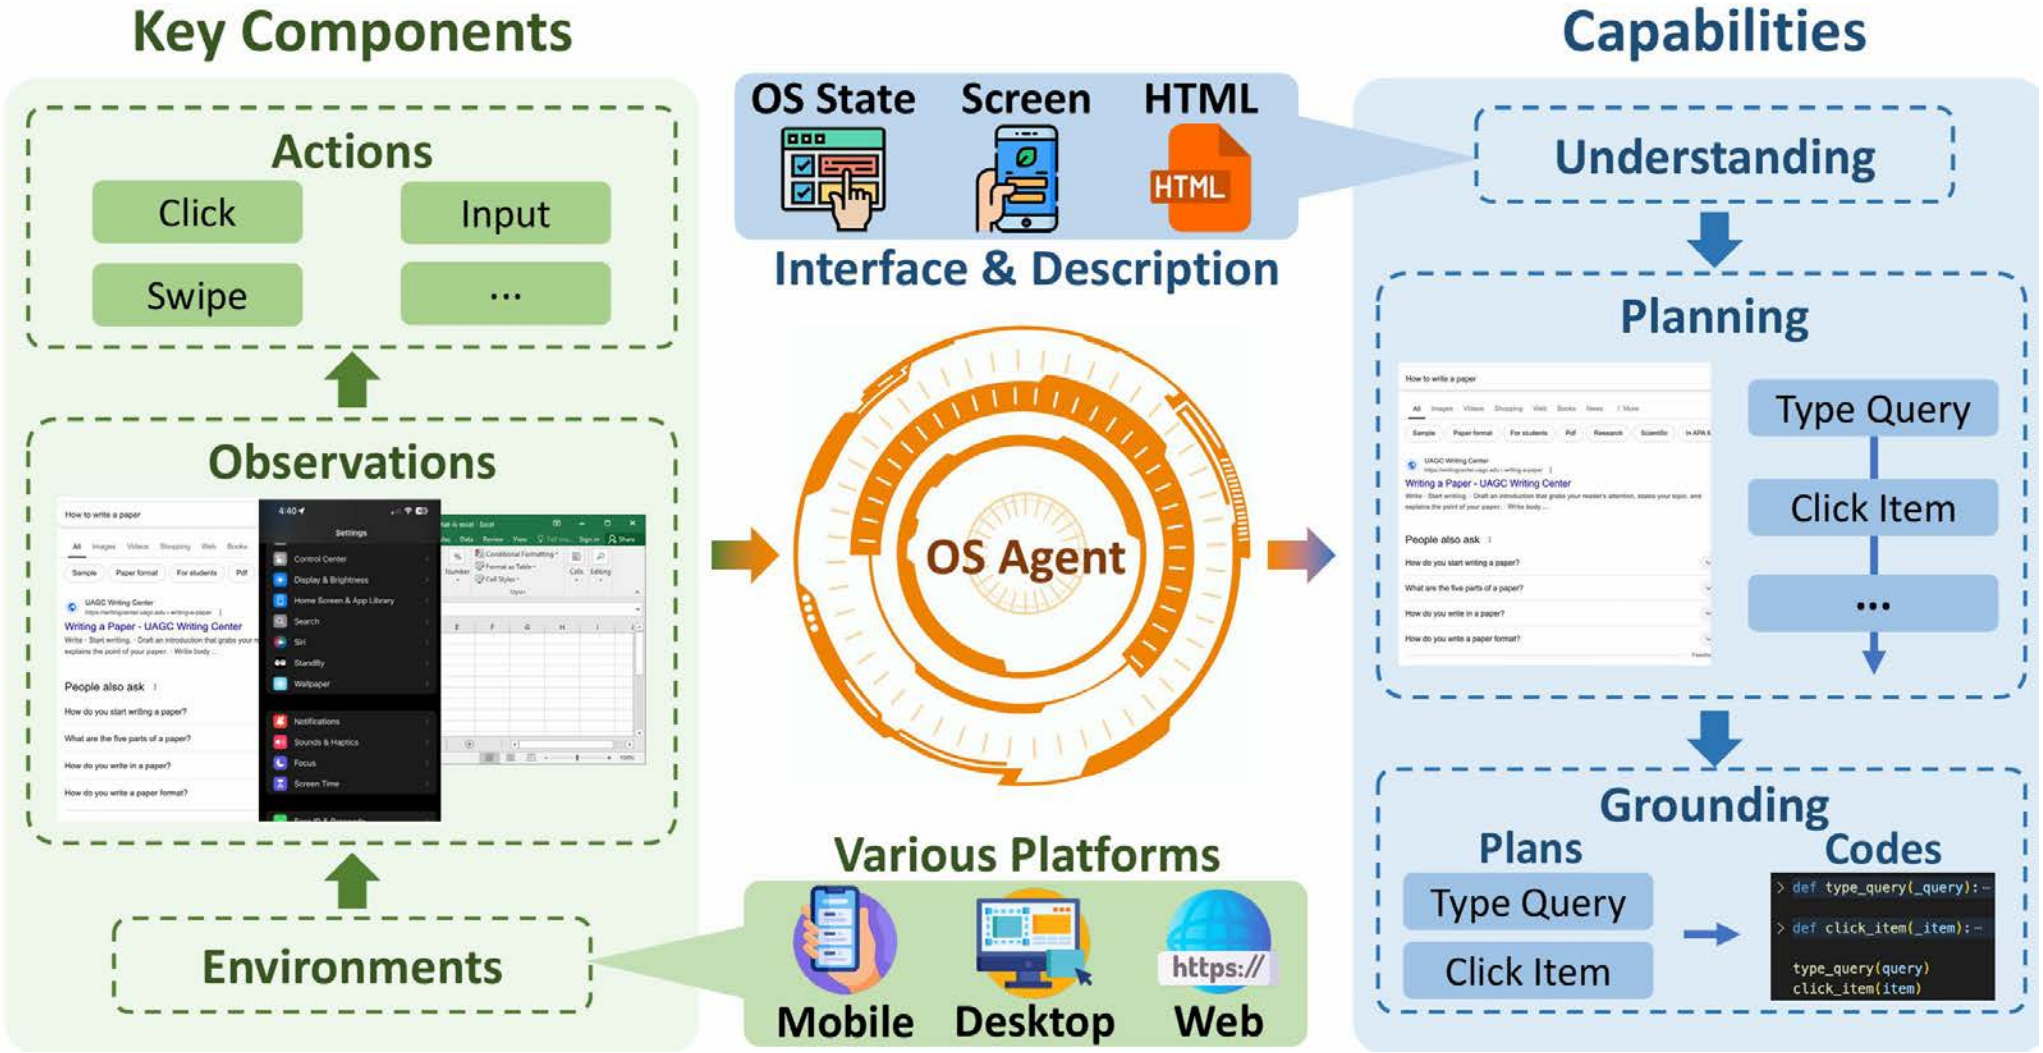

From OS Agents

Agent Application: Computer-using Agent

- Operator

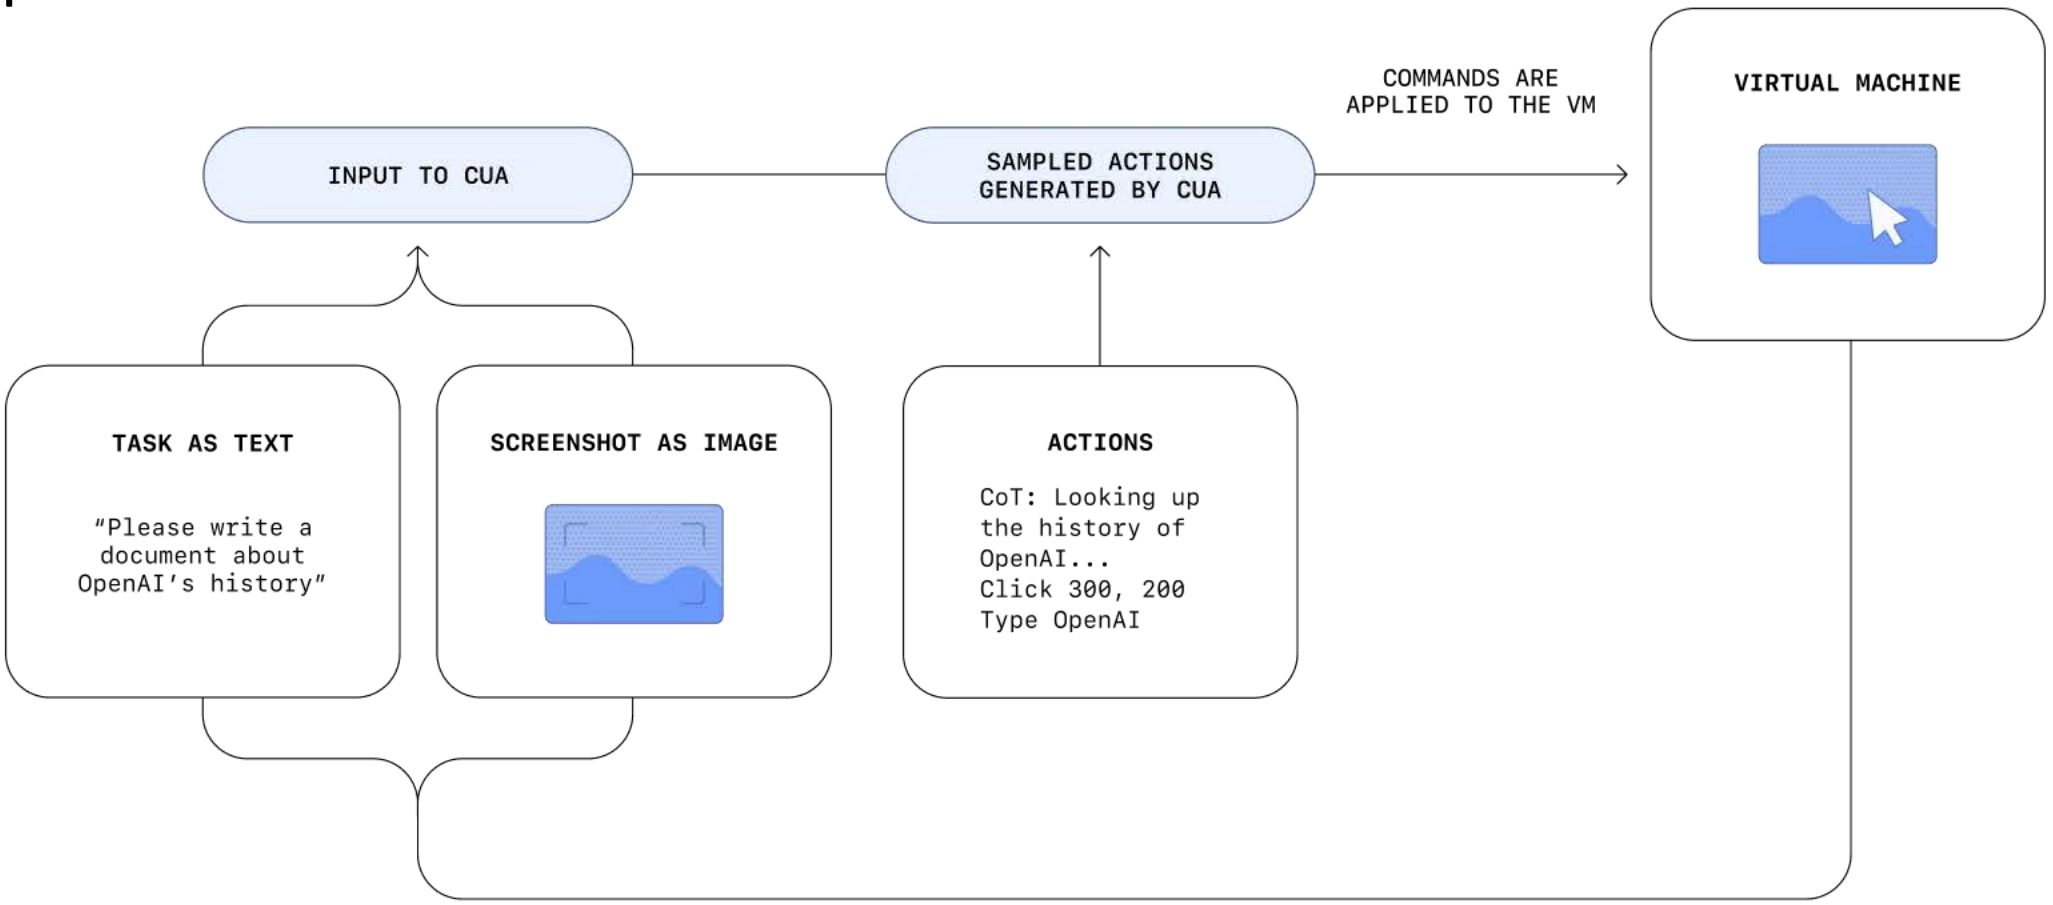

From [Introducing Operator | OpenAI]

Agent Application: Computer-using Agent

- Manus
  - Operator + DeepResearch

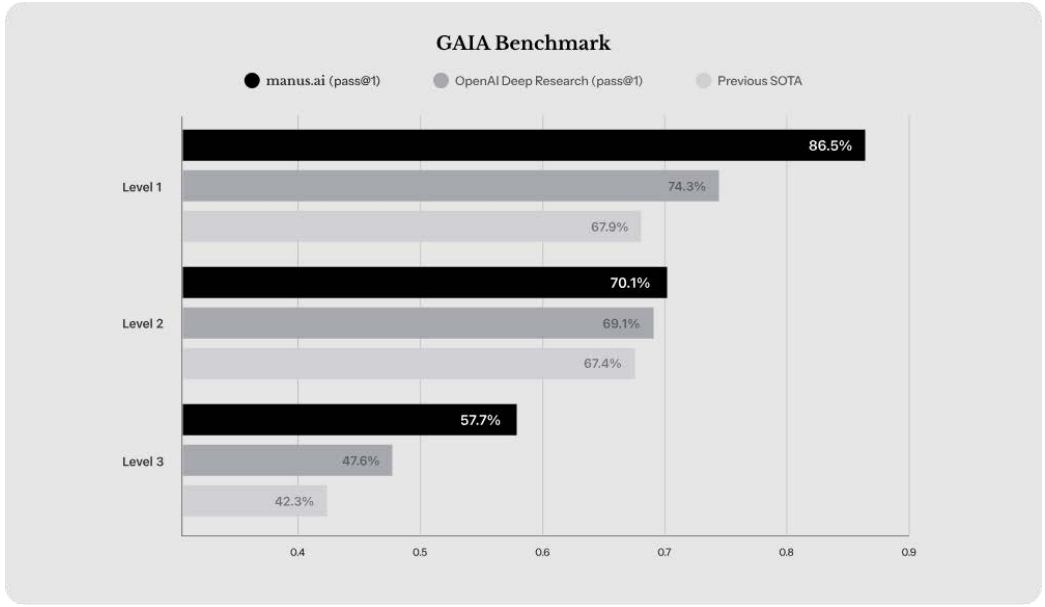

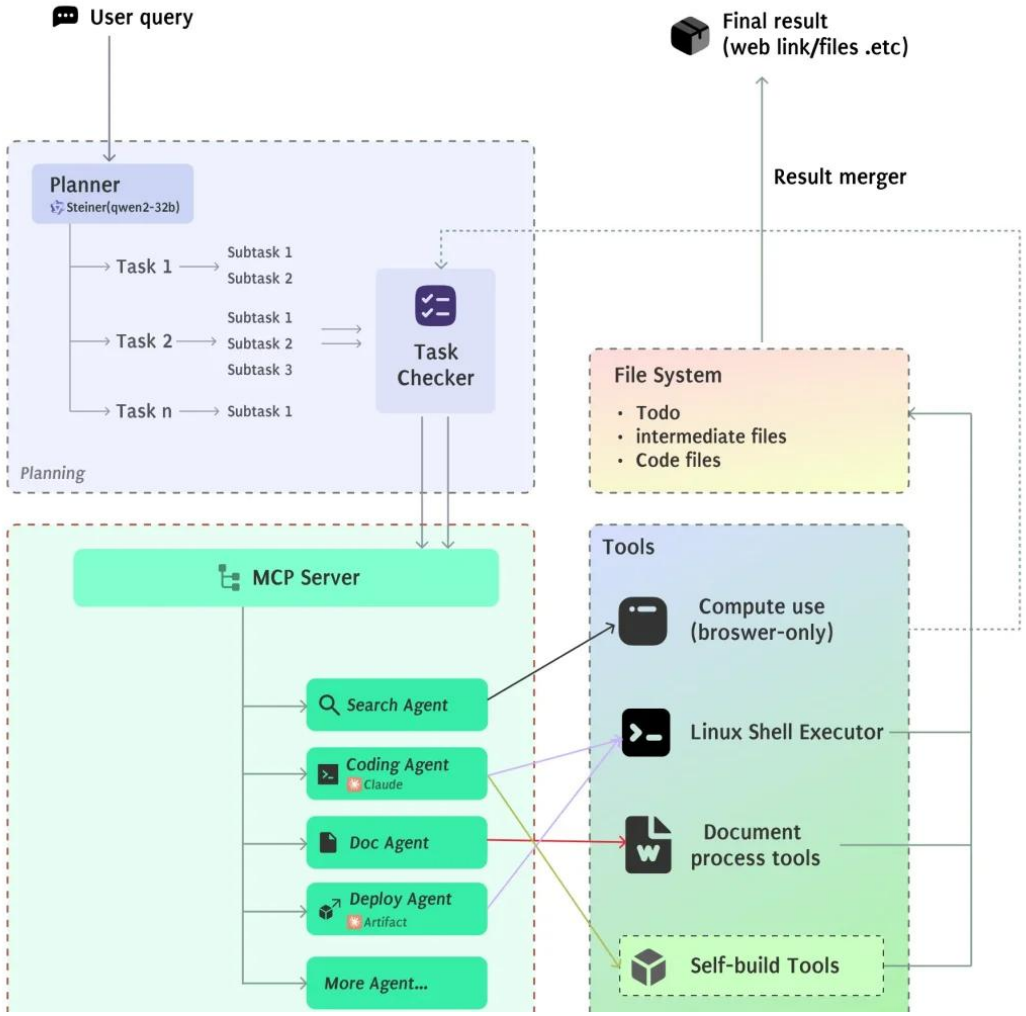

Agent Application: Vertical

- Coding & Development
  - Devin, Cursor, Replit, Windsurf, Trae
- Customer service and Sales
  - Decagon, Clay
- Business research
  - Hebbia
- Scientific research
  - Elicit
- Supply chain
  - Palantir
- Healthcare
  - Epic

Agent Evaluation

- Benchmark
  - OSWorld
  - WebArena
  - WebVoyoger
  - GAIA
  - Humanity’s last exam

Outline

- Background
  - Agent & AGI
  - What is Agent
- Agent: Foundations
  - Key Components: Reasoning, Memory, Tool Use
  - Agentic Workflow vs Large Agent Model
- Agent: Applications
  - Search/Research Agent
  - Computer-Using Agent
  - Other Vertical Agents
- Challenges & Future Trends

Agent Moore Law

- Agentic time double every 7 months

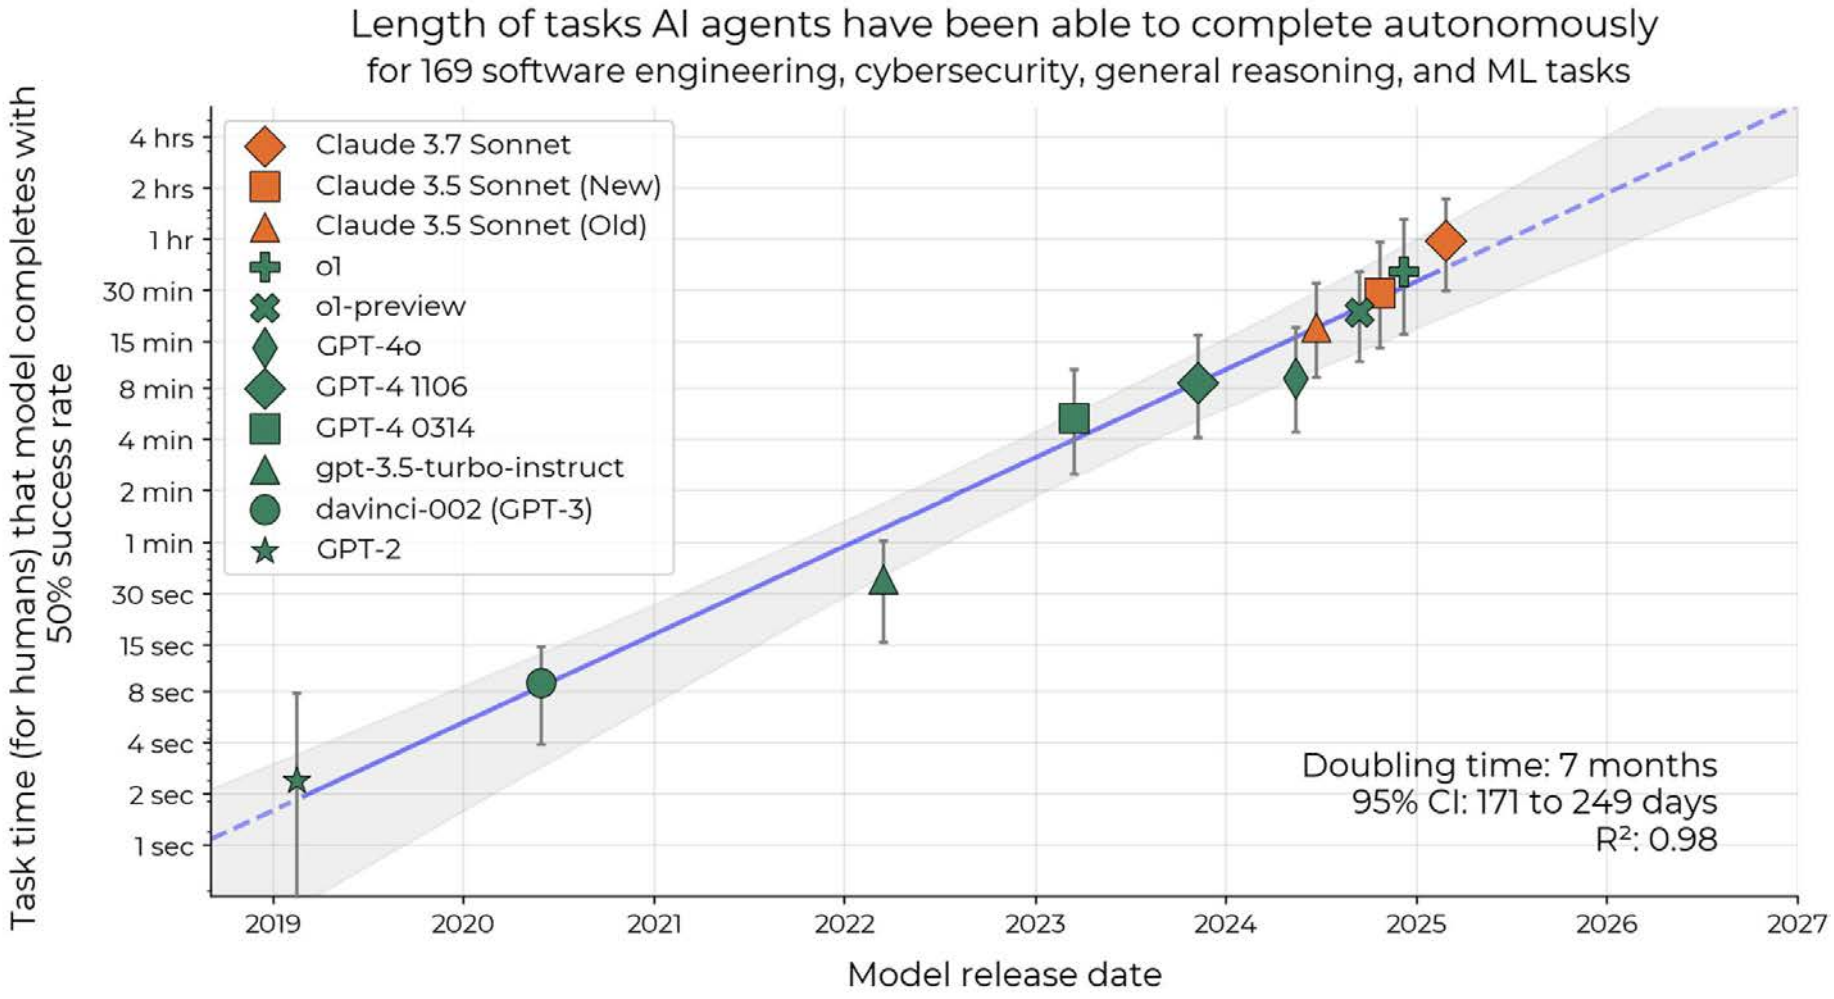

From 2503.14499

Multi-Agent System

- An example

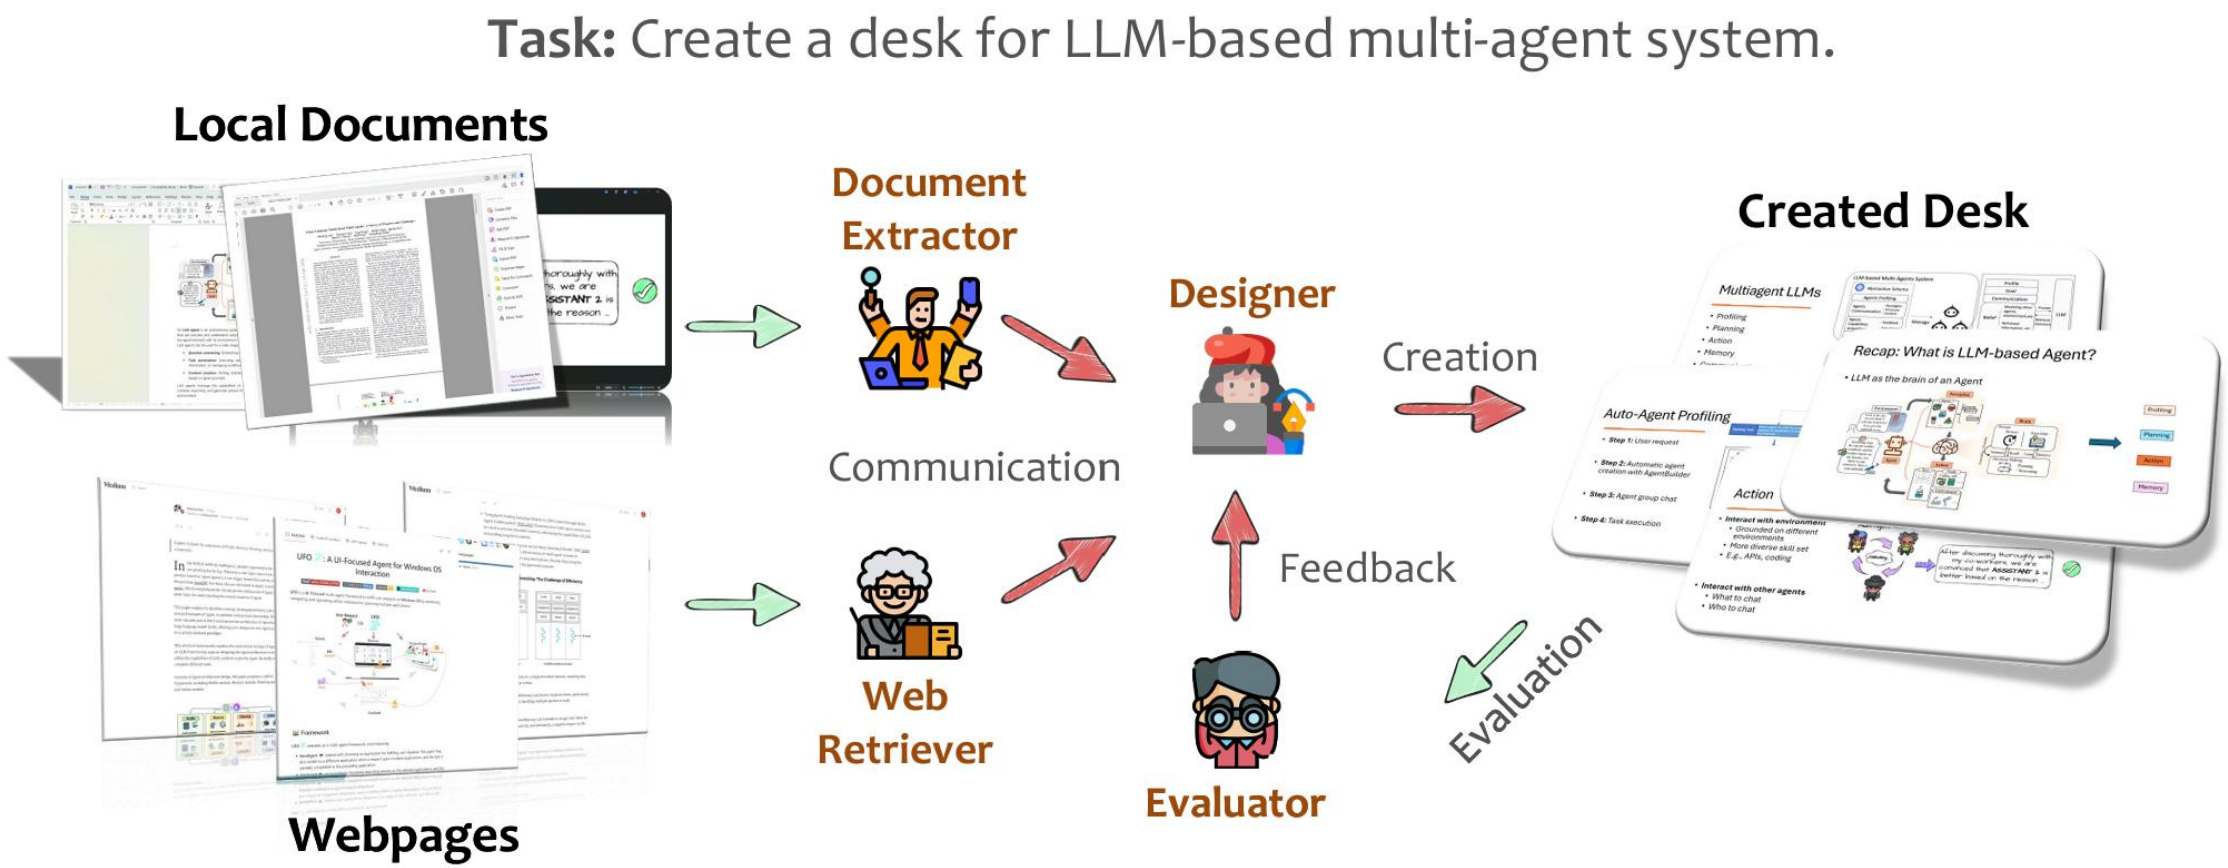

From 2411.18279

Multi-Agent System

- Communication in multi-agent system

Multi-Agent System

- Why multi-agent system fails?
  - Planning: task decomposition, roll assignment
  - Coordination
  - Verification

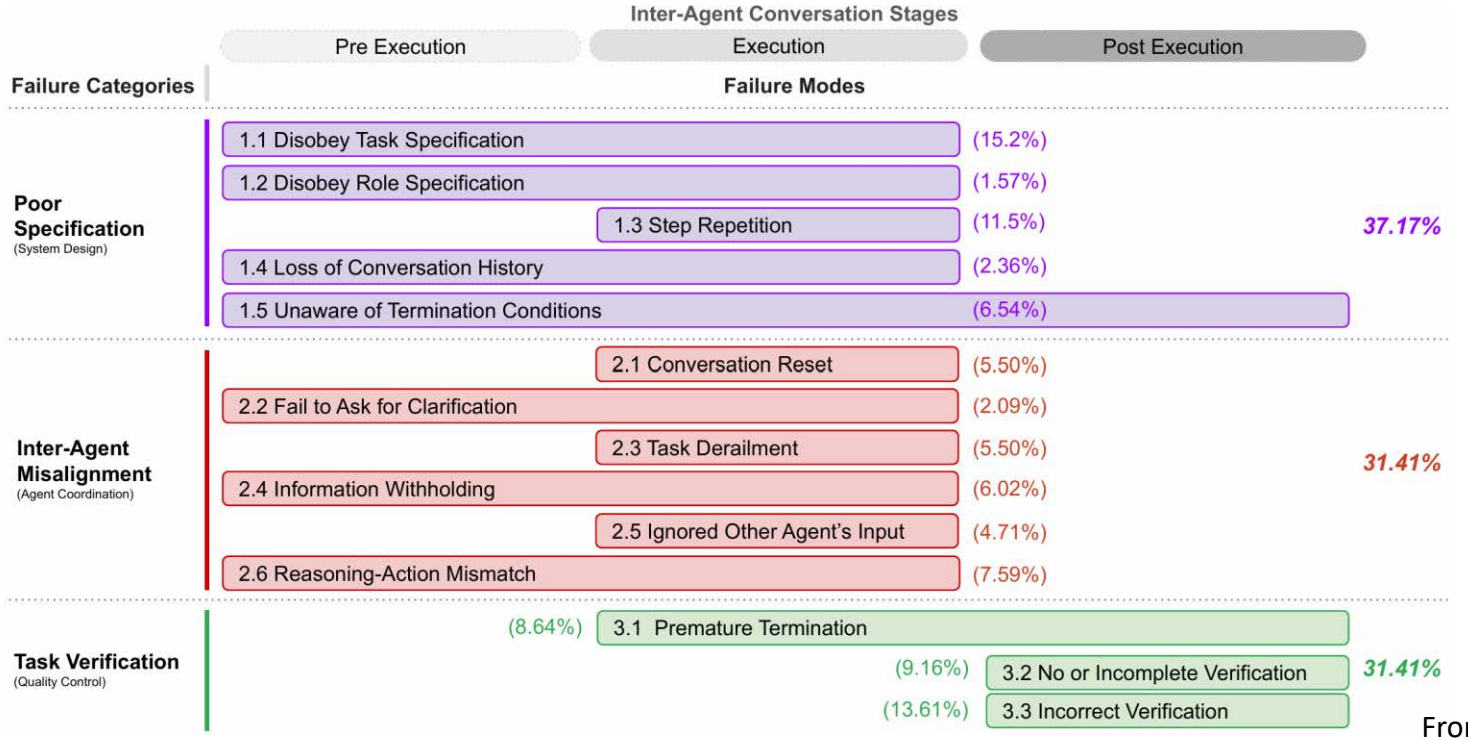

From [Why Do Multi-Agent LLM Systems Fail?]

The Bitter Lesson for Agent

- Less structure, more intelligence

Sutton 这样总结道：

> 「我们必须学会苦涩的教训：人为地去预设我们思考的方式，长期来看并不奏效。AI 研究的历史已经反复验证：
>
> 1）研究者经常试图将知识提前写入智能体；
>
> 2）这种做法短期内效果明显，也让研究者本人很有成就感；
>
> 3）但长期来看，性能很快达到上限，甚至阻碍后续发展；
>
> 4）最终的突破反而来自完全相反的方法，即通过大量计算资源进行搜索和学习。最终的成功让人有些苦涩，因为它否定了人们偏爱的、以人为中心的方法。」

Current Agent Research Ecosystem

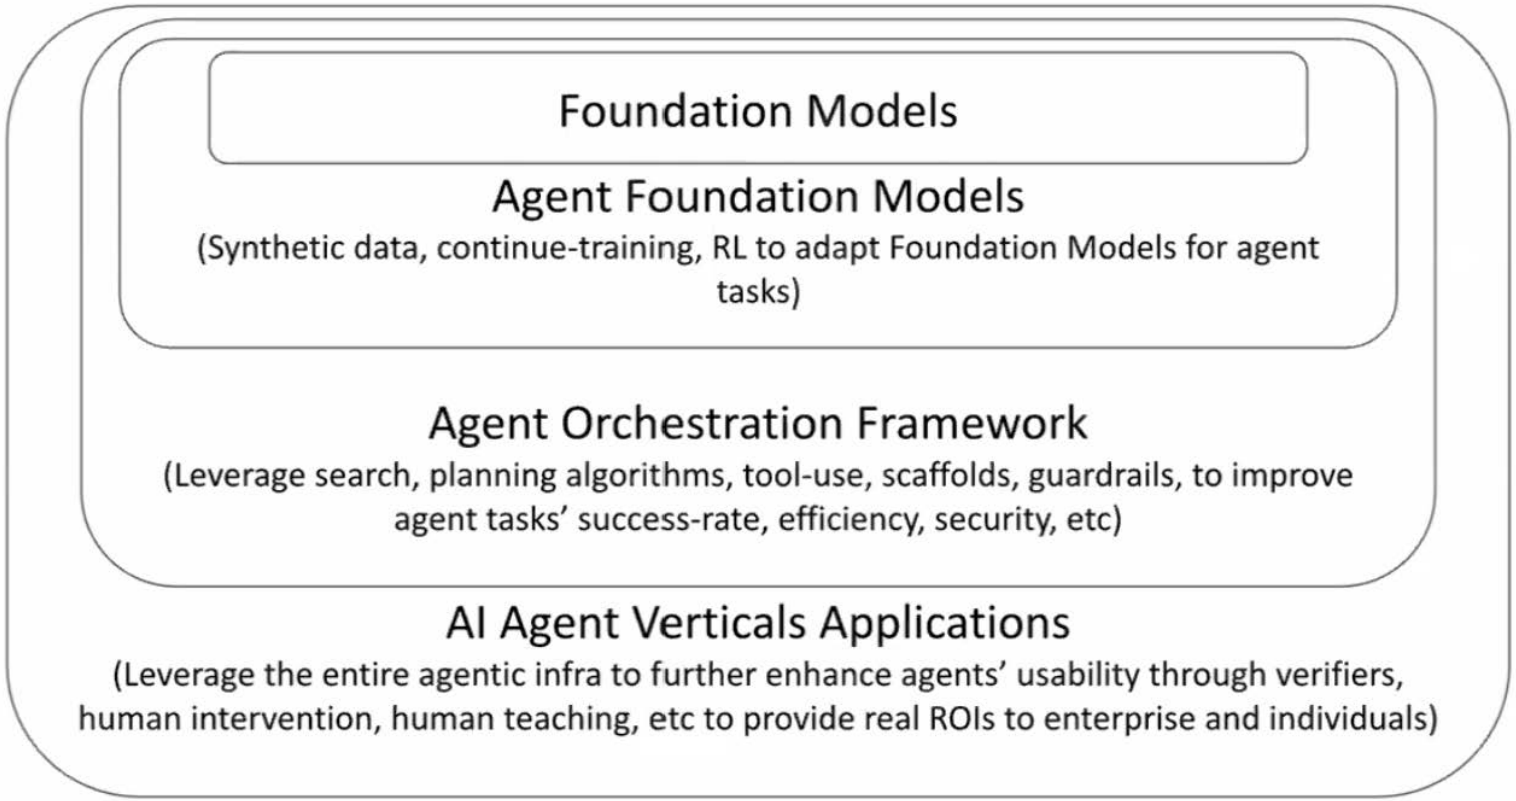

Internalize Reasoning/Action into LLMs

- How to internalize
  - Environment
  - Data
  - Feedback loop
  - Reward model

AGI Roadmap: From Chatbot to Reasoner to Agent

Level 1: Chatbot
(Language model)

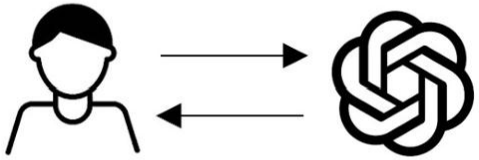

direct respond

Level 2: Reasoner
(Reasoning model)

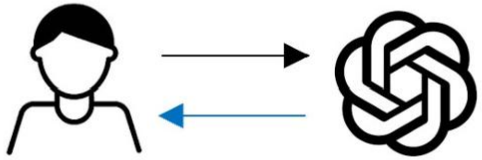

slow think
before respond

Level 3: Agent
(Agent model)

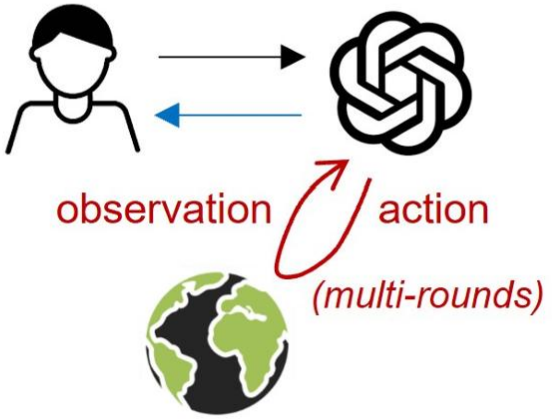

iterative slow think & action
before respond

From 2503.06580

AGI Roadmap: From Agent to Innovator to Organization

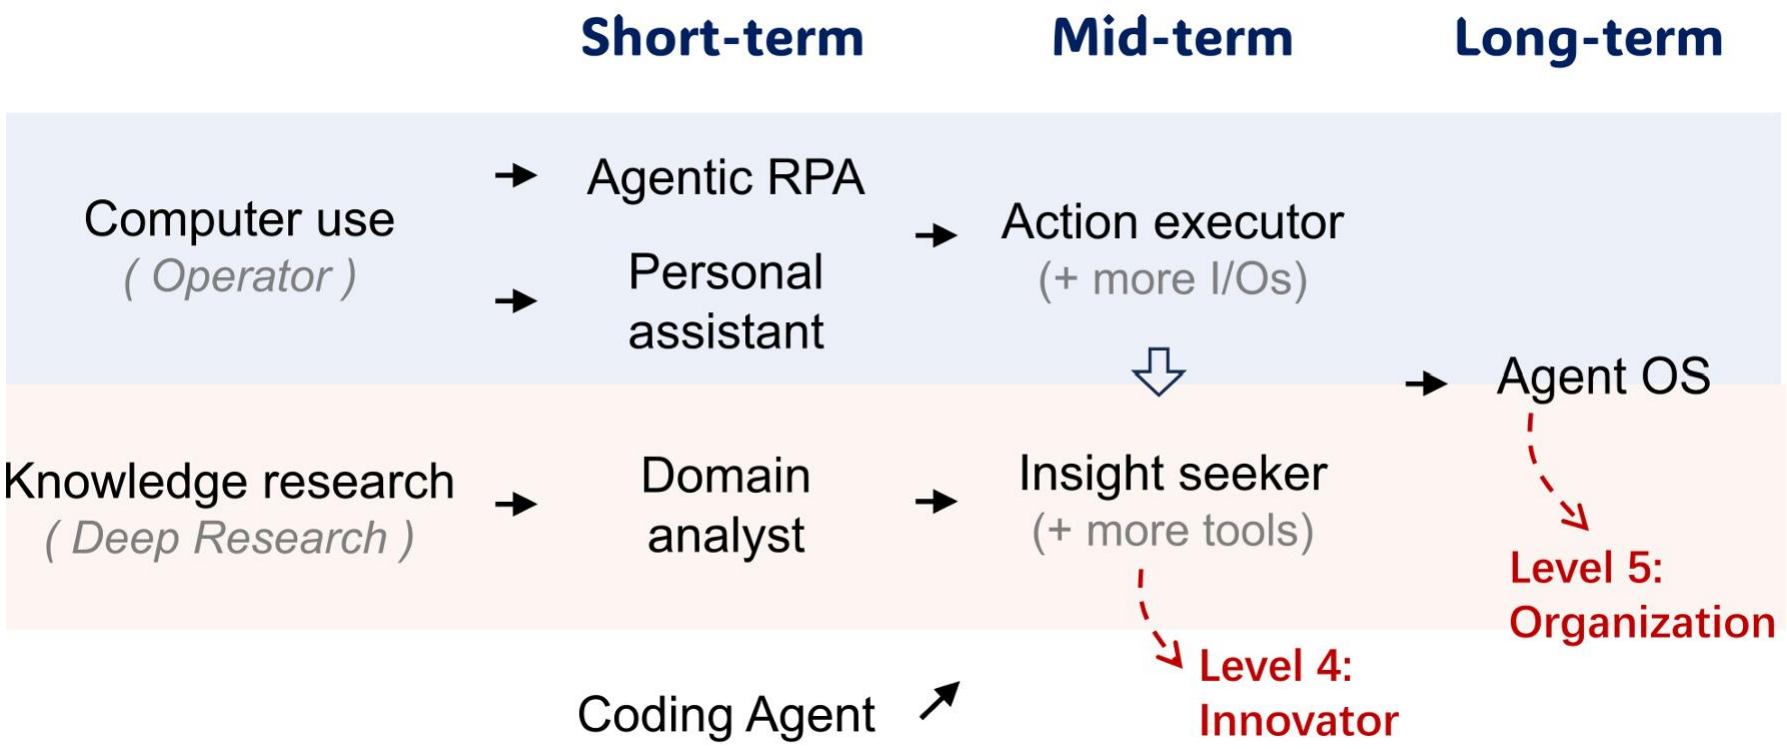

From 2503.06580

AGI Roadmap


OpenAI's 5-Level Plan to AGI
Conversational AI
1
Reasoners
2
Agents
3
Innovators
4
Organizations
5
chatgpt4online.org

Thanks

Xu Tan/谭旭
tanxu2012@gmail.com